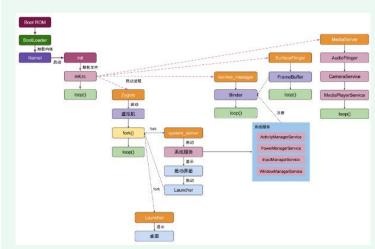


Surfaceflinger





Explain all the details in simple words

This diagram explains the **Android boot process** from powering on the device to launching the home screen (Launcher). Let's walk through it **step-by-step in simple terms**:

1. Power-On and Bootloader Stage

- Boot ROM
 - ➤ This is built into the hardware. It starts automatically when you power on the phone.
- BootLoader
 - ➤ Loads and starts the Android **Linux Kernel** into memory.
- Kernel
 - ➤ The heart of the operating system. It sets up basic things like memory and hardware.

2. Init Process and Starting Services

- Init
 - ➤ The first Android program started by the kernel.
 - > Reads the init.rc file to know what to do next.
- init.rc
 - ➤ A script that tells init what services to start, like:
 - Zygote
 - SurfaceFlinger

- MediaServer
- service_manager

🧵 3. Core Processes Start

- a) service_manager
- ➤ Manages and registers Android system services (like ActivityManager, PowerManager).
- b) SurfaceFlinger
- ➤ Manages the screen display how apps/layers are shown.
- c) MediaServer
- ➤ Handles audio/video playback and camera services.
- d) Binder
- ➤ Enables different Android components to talk to each other.

५ 4. Zygote and Forking

- Zygote
 - ➤ A special process that loads the Android framework and Java libraries.
 - ➤ It is like a "template" used to quickly create other apps or system services.
- fork()
 - ➤ A way to make a copy of Zygote when needed:
 - For **system_server** (which runs core system services)
 - For apps like **Launcher**

5. system_server and System Services

- system_server
 - ➤ Starts essential **system services**, such as:
 - **ActivityManagerService** → Handles activities (UI screens).
 - **PowerManagerService** → Manages power and sleep.
 - **InputManagerService** → Handles touches and key presses.
 - **WindowManagerService** → Handles window positioning and focus.



6. Launcher (Home Screen)

- Zygote forks a new process to start the **Launcher app**.
- The Launcher is the first visible app you see (home screen).
- It allows the user to open apps and interact with the phone.

loop()

Wherever you see loop():

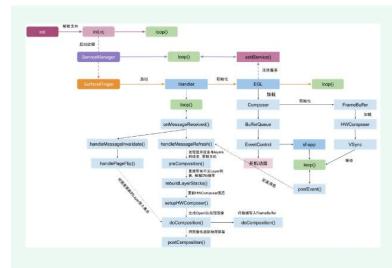
➤ It means the process enters a **main loop**, waiting for messages/events to handle — this is how Android apps and services stay alive and responsive.

⊗ Summary (Flow)

- 1. Boot ROM \rightarrow Bootloader \rightarrow Kernel
- 2. Kernel \rightarrow init \rightarrow init.rc
- 3. init.rc → Starts: Zygote, service_manager, MediaServer, SurfaceFlinger
- 4. Zygote → fork system_server → start system services
- 5. Zygote \rightarrow fork Launcher \rightarrow show home screen

Let me know if you'd like this explained with real file names or logs (like logcat) from an actual Android device.





This second diagram focuses on **SurfaceFlinger**, the core Android system service responsible for **compositing and rendering the screen**. Let's break it down step-by-step in **simple words**, just like a



1. Boot Sequence (relevant parts)

From init.rc:

- init.rc tells the system to start SurfaceFlinger.
- It also starts **ServiceManager** and other services.
- These services run in a loop waiting for commands.

2. SurfaceFlinger Starts

SurfaceFlinger Initialization:

- SurfaceFlinger starts and:
 - Registers itself with ServiceManager using addService() so other services/apps can find it.
 - Sets up **EGL**, which is a library that helps SurfaceFlinger talk to **OpenGL ES or Vulkan** to render things.
 - Starts its own **loop()** to wait for screen update messages.



3. EGL & Rendering Setup

EGL Setup:

- Initializes rendering components:
 - **Composer**: Coordinates what should be drawn.
 - **BufferQueue**: Queue of images (layers) to be rendered.
 - **EventControl**: Manages animation timing (like boot animation).
 - FrameBuffer: Actual screen memory to write to.
 - **HWComposer**: Hardware abstraction for composition (blending images together).



🞡 4. Frame Update Flow — Handling Messages

SurfaceFlinger's Event Loop

• SurfaceFlinger runs a **Handler** loop that waits for events like:



- Something changed and needs to be redrawn (invalidate).
- It's time for a refresh (based on VSYNC signal).

5. Message Handlers — How Drawing Happens

a) handleMessageInvalidate()

- Something changed, needs redraw.
- Calls handlePageFlip() to prepare next frame.

b) handleMessageRefresh()

- Called regularly (usually every 16ms for 60fps).
- Triggers a full frame redraw:
 - 1. **preComposition()** Prepare layer list.
 - 2. **rebuildLayerStacks()** Figure out which layers are visible and how.
 - 3. **setupHWComposer()** Tell hardware composer how to blend them.
 - 4. **doComposition()** Actually draw the layers.
 - If GPU is used → render OpenGL image.
 - If hardware composer is used → let it blend layers directly.
 - 5. **postComposition()** Display the result on screen.

🔁 6. VSync, Animation, and Message Loop

VSync (Vertical Sync):

- Comes from display hardware.
- Tells SurfaceFlinger "time to draw a new frame".

sf-app (SurfaceFlinger animation app):

- A small part that handles animations (like boot animation).
- Uses postEvent() to send signals back into SurfaceFlinger's loop.

🚅 Summary – How the Screen Gets Updated

1. init.rc starts SurfaceFlinger.



- 2. SurfaceFlinger sets up rendering (EGL, BufferQueue, FrameBuffer).
- 3. VSync triggers frame update events.
- 4. SurfaceFlinger's **Handler** receives messages:
 - If something changed → handleMessageInvalidate()
 - On every VSync → handleMessageRefresh()
- 5. Refresh flow:
 - Compose layers → Draw using GPU or hardware → Output to screen.

🖈 Real-World Analogy

Think of SurfaceFlinger like a **movie director**:

- It gets **frames (layers)** from apps.
- It decides how to layer them.
- Then it **asks the camera (GPU/HWC)** to shoot the final scene.
- Finally, it **displays** the shot on the screen.

Let me know if you'd like this diagram tied to real source files or log tags (like **SurfaceFlinger**, **BufferQueueProducer**, etc.).



```
... cc_binary {
  name: "surfaceflinger",
  defaults: [ "surfaceflinger_defaults" ],
  init_rc: [ "surfaceflinger.rc" ],
  srcs: [ "main_surfaceflinger.cpp" ],
  whole_static_libs: [
    "libsigchain",
  ],
  shared libs: [
    "android.frameworks.displayservice@1.0",
     "android.hardware.configstore-utils",
     "android.hardware.configstore@1.0",
     "android.hardware.graphics.allocator@2.0",
     "libbinder",
     "libcutils",
     "libdisplayservicehidl",
     "libhidlbase",
     "libhidltransport",
```



```
"liblayers_proto",
     "liblog",
     "libsurfaceflinger",
     "libtimestats_proto",
     "libutils",
  ],
  static_libs: [
     "libserviceutils",
     "libtrace_proto",
  ],
  Idflags: [ "-WI,--export-dynamic" ],
  // TODO(b/71715793): These version-scripts are required due to the use of
   // whole_static_libs to pull in libsigchain. To work, the files had to be
   // locally duplicated from their original location
   // $ANDROID_ROOT/art/sigchainlib/
  multilib: {
    lib32: {
       version_script: "version-script32.txt",
    },
    lib64: {
       version_script: "version-script64.txt",
    },
  },
} ...
```

This cc_binary build rule describes how the **SurfaceFlinger binary** is built in Android using Soong (the Android.bp build system).

Let's explain each part in simple terms:

```
cc_binary { ... }
```

Defines a **native** (C/C++) **binary** that will be compiled — in this case, it's for SurfaceFlinger.

Basic Properties

| Field | Meaning |
|--|--|
| name: "surfaceflinger" | The name of the final binary. Installed as /system/bin/surfaceflinger. |
| <pre>defaults: ["surfaceflinger_defaults"]</pre> | Inherits common settings (like CFLAGS, include paths) from a default module. |
| <pre>init_rc: ["surfaceflinger.rc"]</pre> | Installs this init script so init can start SurfaceFlinger during boot. |



| Field | Meaning |
|--|---|
| <pre>srcs: ["main_surfaceflinger.cpp"]</pre> | This is the main source file that contains main() for SurfaceFlinger. |

Static Libraries

| Field | Meaning |
|---|--|
| <pre>whole_static_libs: ["libsigchain"]</pre> | Links in all object files from this static library. libsigchain helps handle signal chaining (for e.g., signal handling in ART). |
| static_libs: [] | Statically links in specific utility libraries (like libtrace_proto for tracing). These are linked once per usage. |

Shared Libraries

These are **dynamically linked at runtime** and represent shared system services or utility libraries:

- AIDL/HIDL Interfaces (hardware or system services):
 - android.frameworks.displayservice@1.0
 - android.hardware.configstore@1.0
 - android.hardware.graphics.allocator@2.0
 - libdisplayservicehidl
- System IPC & Utils:
 - libbinder → Binder IPC
 - libhidlbase, libhidltransport → HIDL support
 - libutils, libcutils → Android base utilities
 - liblog → Logging
 - libsurfaceflinger → Main logic for SurfaceFlinger
 - liblayers_proto, libtimestats_proto → Protobuf definitions

These allow SurfaceFlinger to **interact with other system components**, allocate graphics buffers, and talk to hardware.

Linker Flags

| Field | Meaning |
|---|---|
| <pre>ldflags: ["-Wl,export- dynamic"]</pre> | Tells the linker to export all symbols for dynamic linking or for tools like backtraces/profilers. |

🔀 multilib & Version Scripts

| Field | Meaning | |
|----------------|---|--|
| multilib | Sets separate configuration for 32-bit (lib32) and 64-bit (lib64) builds. | |
| version_script | These control which symbols are exported. Needed because libsigchain is statically linked with whole_static_libs, which can cause symbol conflicts. | |

The comment mentions that these version scripts were copied locally from art/sigchainlib/.

Final Output

This builds:

- A binary named surfaceflinger
- Includes HIDL, Binder, and rendering logic
- Uses an init script surfaceflinger.rc to start the process during boot
- Installed to the system partition and runs as part of Android's graphics stack

Let me know if you'd like to:

- Trace the startup from main_surfaceflinger.cpp
- See what's inside surfaceflinger.rc
- Understand how libsurfaceflinger is implemented



SurfaceFlinger Code Analysis [Android 11]

Code path: /frameworks/native/services/surfaceflinger/

The SurfaceFlinger binary is divided into the surfaceflinger executable file (main entry) and the libsurfaceflinger.so library file (function implementation), which is compiled from the main_surfaceflinger.cpp file. The Android.bp code module compilation configuration is as follows:

1. [Executable file - surfaceflinger]

You're absolutely right — in **Android 11**, the **SurfaceFlinger** binary is structured into two parts:

1. SurfaceFlinger Code Structure Overview



• Path:

/frameworks/native/services/surfaceflinger/

- 🧈 Components:
 - main_surfaceflinger.cpp → This is the entry point (main() function).
 - libsurfaceflinger.so → Contains most of the core logic, such as rendering, composition, and interaction with HWC (Hardware Composer).

2. Compilation Overview – Android.bp

You already posted the cc_binary module for surfaceflinger. Let's summarize its role:

A. main_surfaceflinger.cpp (Executable Entry Point)

Type: cc_binary

What it does:

- Starts the SurfaceFlinger service by:
 - Setting up threads
 - Creating the **SurfaceFlinger** class object
 - Initializing EGL/GL and related graphics components
 - Registering the service via Binder
 - Entering the main event loop

Code snippet from main_surfaceflinger.cpp:

```
int main(int, char**) {
    ...
    sp<SurfaceFlinger> flinger = new SurfaceFlinger();
    ...
    flinger->init();
    ...
    defaultServiceManager()->addService(...);
    ...
    IPCThreadState::self()->joinThreadPool();
}
```

B. libsurfaceflinger (Shared Library)

Type: cc_library_shared

What it contains:



- The actual SurfaceFlinger class implementation.
- Handles:
 - Layer management
 - Composition strategy
 - HWComposer integration
 - BufferQueue handling
 - VSync
 - Transactions
 - Frame rendering pipeline

Key source files:

- **SurfaceFlinger.cpp** the core class implementation
- Layer.cpp / LayerBE.cpp represent compositable units (each app window or surface is a Layer)
- RenderEngine/ abstraction for OpenGL/Vulkan rendering
- Scheduler/ manages vsync and frame timing
- DisplayHardware/ abstracts physical display and vsync driver
- CompositionEngine/ coordinates layer composition, especially post Android 10/11 modularization

Relationship Between Executable and Library

| Executable (surfaceflinger) | Shared Library (libsurfaceflinger.so) |
|------------------------------|---|
| Defines main() | Defines the SurfaceFlinger class |
| Calls SurfaceFlinger::init() | Implements rendering logic |
| Registers service | Implements service behavior |

🌉 Modular Design Benefit

By splitting the entry (main_surfaceflinger.cpp) and implementation (libsurfaceflinger), Android:

- Allows other modules to reuse libsurfaceflinger
- Makes testing, maintenance, and modularization easier
- Keeps the main() logic minimal and focused

Would you like a full call flow diagram for **SurfaceFlinger::init()** and how it gets to composition (e.g., VSync \rightarrow invalidate \rightarrow refresh \rightarrow draw)?





cc_library_shared,cc_binary why we need both

Great question!

In Android's build system, using **both** cc_library_shared and cc_binary for something like SurfaceFlinger helps **separate concerns** between:

- the executable entry point (cc_binary)
- the reusable logic and implementation (cc_library_shared)

Let's break it down clearly:

Why we need both cc_binary and cc_library_shared

| Purpose | cc_binary | cc_library_shared |
|-------------------|--|--|
| What it builds | A standalone executable (e.g., surfaceflinger) | A shared library (libsurfaceflinger.so) |
| Contains | Just main() and service registration logic | All the functional logic of SurfaceFlinger |
| Used by | Boot/init system (calls the binary) | Binary and possibly other modules or services |
| Reusable? | X No − it's an app or service entry point | ✓ Yes – other processes can dynamically link to it |

🧱 Real Example: SurfaceFlinger

cc_binary → surfaceflinger

```
cc_binary {
    name: "surfaceflinger",
    srcs: ["main_surfaceflinger.cpp"],
    shared_libs: ["libsurfaceflinger", ...]
}
```

- Contains the main() function.
- Instantiates SurfaceFlinger class.
- Starts Binder service and joins the main loop.

$cc_library_shared \rightarrow libsurfaceflinger.so$

```
cc_library_shared {
   name: "libsurfaceflinger",
   srcs: ["SurfaceFlinger.cpp", "Layer.cpp", ...]
}
```

- Implements all the functionality.
- Provides SurfaceFlinger::init(), SurfaceFlinger::handleMessageRefresh(), etc.
- Can be reused/tested independently.

🔁 Analogy

Think of cc_library_shared as a car engine, and cc_binary as the car:

- You build the **engine separately** so it can be reused in other cars (or tested by itself).
- The **car binary** (main()) knows how to **start the engine**, register with traffic (Binder), and drive (event loop).

Benefits of Splitting

| Benefit | Description |
|-------------|--|
| Reusability | Other services (like SurfaceFlinger_test , or tools like dumpstate) can link against libsurfaceflinger . |
| Testability | You can unit test libsurfaceflinger functions without needing to launch the whole binary. |
| | Keeps the binary small and clean. Complex code stays in the shared library. |
| Maintenance | If logic changes, you often only update the . so , not the full binary. |

TL;DR

- cc_binary: Runs the process (has main()).
- cc_library_shared: Implements the logic (reused by binary and others).
- We use **both** so logic is cleanly separated from process entry, reused easily, and maintained modularly.

Let me know if you want a breakdown of the call flow starting from $main() \rightarrow SurfaceFlinger::init() \rightarrow frame drawing!$





For what purpose will others use libsurfaceflinger.so

Excellent question — and a key one for understanding Android's graphics architecture.

Why Others Might Use libsurfaceflinger.so

While surfaceflinger binary is the actual service that runs and handles display composition, the libsurfaceflinger.so shared library exposes core logic and utilities that other parts of the system may **link to**, including:

1. Testing and Debugging Tools

- SurfaceFlinger test binaries (like SurfaceFlinger_test):
 - Link to libsurfaceflinger.so to test internal functions.
 - Useful for unit testing composition logic without launching the full system.

2. System UI or Shell Components

- Android's Shell, SystemUI, or WindowManager may:
 - Use libsurfaceflinger symbols to query or manipulate surface or display state.
 - This is rare and tightly controlled (mostly indirect via Binder IPC), but may still link certain internal headers for debugging or analysis.

3. Debugging/Profiling Tools

- Tools like dumpsys, SurfaceFlinger::dump():
 - The binary uses public dump() APIs implemented in libsurfaceflinger.so to collect surface, layer, and composition info.
 - These symbols are reused for command-line and UI-level debugging.

Example:

bash

adb shell dumpsys SurfaceFlinger



4. Binder IPC Glue Code

Even though most components talk to SurfaceFlinger via Binder IPC, the IPC stub/proxy code may depend on types or utility functions provided by libsurfaceflinger.

- For example:
 - ISurfaceComposer or DisplayDevice types may be declared in shared headers linked with libsurfaceflinger.

5. Layer Management Utilities

If some system service or tool (e.g. dumpsys, simpleperf, SurfaceTracing) wants to:

- Inspect layer trees
- Analyze frame rendering
- Log time spent per layer Then it may reuse libsurfaceflinger.so to avoid duplicating logic.

6. Vendor Extensions / Debug Builds

In debug or custom vendor builds:

- OEMs (like Samsung, Xiaomi) may extend or patch SurfaceFlinger.
- Instead of modifying the binary, they may dynamically hook or wrap around the .so to:
 - Inject frame capture
 - Customize composition logic
 - Export additional logs

Summary of Consumers of libsurfaceflinger.so

| Module | Why it uses libsurfaceflinger.so |
|----------------------------|-------------------------------------|
| surfaceflinger (binary) | Main service |
| dumpsys | Access layer and frame state |
| SurfaceFlinger_test | Unit testing |
| simpleperf, SurfaceTracing | Performance tracing |
| Vendor tools / OEM patches | Custom diagnostics or frame capture |
| Debug UI / Shell tools | Read internal Surface state |



Note:

Most apps or services **do not directly link libsurfaceflinger.so** — they use **Binder IPC** via **ISurfaceComposer and ISurfaceComposerClient**. But **internal tools and tests** do link the library to reuse shared logic.

Would you like to see a list of functions exposed by **libsurfaceflinger.so** (e.g., from symbol dumps) or how **ISurfaceComposer** is used across modules?



```
int main(int, char **) {
  signal(SIGPIPE, SIG_IGN);
  // Starting from 8.0, Android provides the HIDL mechanism, which standardizes the original
direct JNI->Native->HAL interface call method into a HIDL service/client interaction method.
  // This method standardizes and unifies the calling mechanism of the Android Framework
and HAL. However, from a project perspective, this calling method has a performance overhead
and will take more time than the direct calling method.
  hardware::configureRpcThreadpool(1 /* maxThreads */,
       false /* callerWillJoin */);
  startGraphicsAllocatorService();
  // When SF is launched in its own process, limit the number of
  // binder threads to 4.
  ProcessState::self()->setThreadPoolMaxThreadCount(4 );
  // start the thread pool
  sp<ProcessState> ps(ProcessState::self());
  ps- > startThreadPool();
  // Create a SurfaceFlinger object, pointed to by a strong pointer.
  // SurfaceFlinger inherits the RefBase class, so once the new object is assigned to the sp
pointer, the onFirstRef method of the SurfaceFlinger class will be called immediately.
  // instantiate surfaceflinger
  sp<SurfaceFlinger> flinger = surfaceflinger::createSurfaceFlinger();
  setpriority(PRIO_PROCESS, 0, PRIORITY_URGENT_DISPLAY);
  set_sched_policy( 0 , SP_FOREGROUND);
  // Put most SurfaceFlinger threads in the system-background cpuset
  // Keeps us from unnecessarily using big cores
  // Do this after the binder thread pool init
```



```
if (cpusets_enabled()) set_cpuset_policy(0 , SP_SYSTEM);
  // SurfaceFlinger class formally initialized
  // initialize before clients can connect
  flinger-> init();
  // SurfaceFlinger registers the Binder service with ServiceManager,
  // so that other processes can get the SurfaceFlinger service through
getService+SERVICE_NAME, and then can communicate with the SurfaceFlinger class through
Binder.
  // publish surface flinger
  sp<IServiceManager> sm(defaultServiceManager());
  sm ->addService(String16(SurfaceFlinger::getServiceName()), flinger, false,
          IServiceManager::DUMP_FLAG_PRIORITY_CRITICAL |
IServiceManager::DUMP_FLAG_PROTO);
  // The new DisplayService() method inside calls the HIDL definition interface
Return<sp<IDisplayEventReceiver >> getEventReceiver() override;
  startDisplayService(); // dependency on SF getting registered above
  if (SurfaceFlinger::setSchedFifo( true ) != NO_ERROR) {
    ALOGW( "Couldn't set to SCHED_FIFO: %s", strerror(errno));
  }
  // SurfaceFlinger class enters the main loop (note that the SurfaceFlinger class does not
inherit the Threads class and does not follow the interface execution order of the Threads class)
  // run surface flinger in this thread
  flinger-> run();
  return 0;
}
```

This is the main() function of the surfaceflinger binary from Android 11, defined in main_surfaceflinger.cpp. It's the entry point for the SurfaceFlinger system service. Let's break it down line by line in simple language with comments on what each part does.

Step-by-step Explanation

```
int main(int, char**) {
```

Entry point. Parameters unused here.



```
cpp
signal(SIGPIPE, SIG_IGN);
```

Ignore **SIGPIPE**, which can occur if a write is done to a closed socket or pipe. Prevents the process from being killed.

```
cpp
hardware::configureRpcThreadpool(1 /* maxThreads */, false /* callerWillJoin */);
```

This sets up the **HIDL Binder threadpool** for handling remote HAL calls.

- 1: Max 1 thread handles incoming HIDL calls.
- false: Main thread doesn't join the HIDL pool (it's separate).

```
cpp
startGraphicsAllocatorService();
```

Starts the **graphics memory allocator HAL service** (Gralloc), which handles buffer allocations shared between SurfaceFlinger and GPU/HWC.

```
cpp
ProcessState::self()->setThreadPoolMaxThreadCount(4);
```

Limits the **Binder IPC threadpool** to 4 threads. This is important because SF is sensitive to scheduling overhead and latency.

```
sp<ProcessState> ps(ProcessState::self());
ps->startThreadPool();
```

Starts the **Binder IPC thread pool** so SurfaceFlinger can receive IPC calls from clients like WindowManager.

```
cpp
sp<SurfaceFlinger> flinger = surfaceflinger::createSurfaceFlinger();
```



Creates the main **SurfaceFlinger object** using a smart pointer **sp<>**.

• Because SurfaceFlinger inherits RefBase, its onFirstRef() method will be called right after creation.

```
cpp
setpriority(PRIO_PROCESS, 0, PRIORITY_URGENT_DISPLAY);
```

Boosts the priority of this process so it gets **high priority scheduling** — important for rendering to happen on time (avoid jank/stutter).

```
set_sched_policy(0, SP_FOREGROUND);
```

Moves this process to the **foreground scheduling group**, giving it more CPU time compared to background processes.

```
if (cpusets_enabled()) set_cpuset_policy(0, SP_SYSTEM);
```

If **CPUsets** (CPU partitioning feature) are enabled:

- Place SF into the system-background cpuset.
- Helps avoid using big CPU cores unnecessarily (power saving).

```
flinger->init();
```

Calls the init() method of the SurfaceFlinger class:

- Initializes EGL, DisplayDevices, HWComposer, Scheduler, etc.
- Must be done before clients connect (e.g., SurfaceComposerClient, DisplayService).

Registers SurfaceFlinger as a Binder service named "SurfaceFlinger":

- Other processes (like WindowManager, SystemUI) get access using getService("SurfaceFlinger").
- They can send commands like createSurface(), setLayerStack(), etc.

```
cpp
startDisplayService(); // dependency on SF getting registered above
```

Starts the **DisplayService**, which uses SF to handle display events.

• Relies on SF already being registered to function properly.

```
if (SurfaceFlinger::setSchedFifo(true) != NO_ERROR) {
    ALOGW("Couldn't set to SCHED_FIFO: %s", strerror(errno));
}
```

Tries to enable **real-time FIFO scheduling** for SurfaceFlinger.

- Helps ensure strict timing and deadline for frame composition.
- Warns if it fails, but continues running.

```
cpp
flinger->run();
```

Enters **SurfaceFlinger's main loop**:

- Waits for VSYNC, buffer updates, composition requests.
- Handles messages using a Handler-based event loop.
- Keeps rendering frames until the system shuts down.

```
cpp
return 0;
```

Exit code (won't be reached under normal operation).

Overall Flow Summary



| Phase | What Happens |
|------------------------|---|
| ∦ Setup | Signal handling, thread pools, HIDL RPC threadpool |
| Initialization | Create SF object, initialize graphics stack |
| ★ Service Registration | Register SF to ServiceManager |
| | Start DisplayService |
| Real-time Scheduling | Try to boost scheduler priority |
| | Enter run() and handle screen composition, messages, vsync |

Would you like a visual diagram showing how this main() connects to SurfaceFlinger::init(), the VSync loop, and the display pipeline?



hardware::configureRpcThreadpool(1 /* maxThreads */, false /* callerWillJoin */); and ProcessState::self()->setThreadPoolMaxThreadCount(4); why these two

Excellent question! These two lines both set up **threadpools**, but for **different Binder IPC systems** in Android.

Let me explain them clearly and simply:



1. hardware::configureRpcThreadpool(...)

срр

hardware::configureRpcThreadpool(1, false);

Purpose:

This sets up the **HIDL Binder RPC threadpool** (used for hardware services).

Used for:

- HIDL HAL services (e.g., Gralloc, Composer HAL)
- SurfaceFlinger may serve as a HIDL client or server

Meaning of arguments:

- 1: Use **1 thread** to handle incoming HIDL method calls.
- false: Main thread will not join this threadpool.



2. ProcessState::self()->setThreadPoolMaxThreadCount(4);

срр

ProcessState::self()->setThreadPoolMaxThreadCount(4);

Purpose:

This sets up the **AIDL Binder threadpool** — used for **Framework <-> System** communication.

Used for:

- Classic Android Binder calls (e.g., SurfaceComposerClient -> SurfaceFlinger)
- SurfaceFlinger is a **Binder service** used by WindowManager, SystemUI, etc.

What it does:

Allows up to 4 threads to handle incoming AIDL Binder IPCs.

This means SurfaceFlinger can serve up to 4 client requests concurrently via Binder.

🔄 Summary Table

| Function | Purpose | For Which IPC? | Who Uses It |
|---|--------------|----------------------|----------------------------------|
| <pre>hardware::configureRpcThreadpool(1, false)</pre> | Sets up HIDL | HIDL | HAL calls (like Gralloc, Display |
| | thread pool | Binder | HAL) |
| setThreadPoolMaxThreadCount(4) | Sets up AIDL | AIDL | Framework clients (like |
| | thread pool | Binder | WindowManager, SystemUI) |

Why both?

SurfaceFlinger talks to:

- HALs via HIDL → Needs HIDL threadpool (configureRpcThreadpool)
- Apps/System via AIDL Binder → Needs Binder threadpool (setThreadPoolMaxThreadCount)

Without both, it wouldn't be able to:



- Allocate buffers or interact with HWComposer HAL (HIDL)
- Serve system-level clients like WindowManager or SurfaceControl (AIDL)

Let me know if you'd like a diagram showing how SurfaceFlinger connects to other components via HIDL vs AIDL!



```
multilib: {
    lib32: {
       version_script: "version-script32.txt" ,
    },
    lib64: {
       version_script: "version-script64.txt" ,
    },
    },
    why this?
```

Great question! Let's break this down simply.

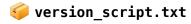
multilib in Android.bp files

```
multilib: {
    lib32: {
        version_script: "version-script32.txt",
    },
    lib64: {
        version_script: "version-script64.txt",
    },
},
```

V Purpose:

To build both 32-bit and 64-bit versions of a native library, each with their own version script.

Why do we need different version scripts?



A version script controls which symbols are exported from a .so file.

This is useful to:



- Hide internal/private functions
- Reduce binary size
- Prevent symbol conflicts
- Maintain ABI stability

Why separate for 32-bit and 64-bit?

Because:

- Symbol layouts may differ between architectures
- You might want to **expose or hide different symbols** on 32-bit vs 64-bit
- For compatibility or performance reasons, the ABI might differ slightly

What happens with this config?

The build system will:

- 1. Compile libsomething.so for 32-bit using version-script32.txt
- 2. Compile libsomething.so for 64-bit using version-script64.txt
- 3. Each version will only expose symbols defined in its respective script

When is this used?

- System libraries that must support both 32-bit and 64-bit apps
- Libraries that are loaded by both 32-bit and 64-bit processes (e.g. libsurfaceflinger.so)
- Ensuring ABI compatibility across both architectures

Summary:

| Field | Meaning |
|----------------------|-----------------------------------|
| multilib | Indicates per-architecture config |
| lib32.version_script | Symbol export list for 32-bit .so |
| lib64.version_script | Symbol export list for 64-bit .so |

This allows fine-grained control over what each version of the library exports.

Let me know if you want to see what's typically inside a version script!



2. [Dynamic library - libsurfaceflinger.so]

The compilation configuration of the Android.bp code module is as follows:

```
... cc_library_shared {
  name: "libsurfaceflinger",
  defaults: [ "libsurfaceflinger_defaults" ],
  cflags: [
    "-fvisibility=hidden",
     "-Werror=format",
     "-DREDUCE_VIDEO_WORKLOAD",
     "-DUSE_AML_HW_ACTIVE_MODE",
  ],
  srcs: [
    ":libsurfaceflinger_sources",
  logtags: [ "EventLog/EventLogTags.logtags" ],
  include_dirs: [
    "frameworks/native/vulkan/vkjson",
     "frameworks/native/vulkan/include",
     "hardware/amlogic/gralloc/amlogic",
     "hardware/amlogic/hwcomposer/tvp",
     "hardware/amlogic/gralloc",
  ],
  static_libs: [
    "libomxutils_static@2",
     "libamgralloc_ext_static@2",
  ],
  cppflags: [
    "-fwhole-program-vtables", // requires ThinLTO
  ],
  Ito: {
    thin: true,
  },
} ....
```

Let's break this cc_library_shared block for libsurfaceflinger.so into simple parts, and explain why each is used.

Overview:

This is the build rule in Android.bp to compile the shared library libsurfaceflinger.so.



This is the core library that implements SurfaceFlinger — Android's compositor service — and is used both internally (by the surfaceflinger binary) and externally (by other system components).



Breakdown of fields:



name: "libsurfaceflinger"

Defines the name of the .so (shared object) library.



```
defaults: [ "libsurfaceflinger_defaults" ]
```

Includes common settings from a default group (likely includes shared srcs, flags, etc.)



```
✓ cflags: [ ... ]
```

Compiler flags for **C/C++ compilation**:

```
-fvisibility=hidden
```

- → Hide all symbols by default (only export those explicitly marked). Helps with:
 - Smaller binary
 - Preventing symbol conflicts

-Werror=format

→ Makes format string warnings into errors (e.g., printf("%d", str) will fail to compile).

```
-DREDUCE_VIDEO_WORKLOAD
```

- -DUSE_AML_HW_ACTIVE_MODE
- → These are **macros** (defines) used in the code to enable Amlogic-specific optimizations, like:
 - Reducing load for video layers
 - Using hardware-specific power-saving or performance modes



srcs: [":libsurfaceflinger_sources"]

Includes source files from a **named source group**. Instead of listing .cpp files here, it reuses another group (likely declared elsewhere in the file).

```
✓ logtags: [ "EventLog/EventLogTags.logtags" ]
```

Used for Android's EventLog. This file defines tags for structured logging in logs.

```
☑ include_dirs:[ ... ]
```

Adds extra **header file search paths** — these are needed for external dependencies:

- Vulkan support: "frameworks/native/vulkan/..."
- Amlogic hardware-specific code: "hardware/amlogic/..."

This shows SurfaceFlinger is integrated with Vulkan and customized for Amlogic chips.

```
✓ static_libs: [ ... ]
```

These are **static libraries** that will be **linked into libsurfaceflinger.so** at build time:

```
text

libomxutils_static@2 → OMX utilities for media
libamgralloc_ext_static@2 → Amlogic Gralloc extension
```

These are architecture-specific helpers for memory and buffer handling.

```
cppflags: [ "-fwhole-program-vtables" ]
```

Enables whole-program vtable optimization (makes C++ virtual function dispatch more efficient). Requires ThinLTO.

```
✓ lto: { thin: true }
```

Enables Thin Link-Time Optimization (LTO):

- Improves performance by optimizing across files
- ThinLTO is faster and uses less memory than full LTO

Summary

| Field | What It Does |
|--------------|--|
| name | Name of the shared library (libsurfaceflinger.so) |
| defaults | Imports common settings |
| cflags | Compiler flags for warnings, symbol visibility, and macros |
| srcs | Source files (via a group) |
| logtags | Tags for event logging |
| include_dirs | Extra header paths (e.g., Vulkan, Amlogic) |
| static_libs | Static .a libraries to link |
| cppflags | Extra C++ optimization flags |
| lto.thin | Enables Thin Link-Time Optimization |

This build config ensures:

- A secure and optimized binary
- Hardware-specific support
- · Compatibility with other system components

Let me know if you want me to explain the .bp for the SurfaceFlinger executable (cc_binary) too!



srcs: [":libsurfaceflinger_sources"]

logtags: ["EventLog/EventLogTags.logtags"]

lto: { thin: true }

explain these in simple words

Here's a simple explanation for each of these Android.bp fields in the cc_library_shared module for libsurfaceflinger:

srcs: [":libsurfaceflinger_sources"]

- What it means: This includes source files (like .cpp) for compilation.
- Why the colon :: It's referencing another **module** named **libsurfaceflinger_sources**, not just raw files.
- That module contains the actual source file list (srcs) used to build libsurfaceflinger.

Simple analogy: It's like saying:



"Get the list of source files defined in the libsurfaceflinger_sources recipe and use them here."

- ✓ logtags: ["EventLog/EventLogTags.logtags"]
 - What it means: It adds Android logtag definitions, which are used to generate EventLog constants.
 - These tags are used when logging system events (e.g., via ALOG_EVENT()).
 - Android will auto-generate a header file (e.g., EventLogTags.h) from this .logtags file.

Why: To map log event IDs to names consistently, like:

```
CDD
LOG_EVENT_INT(MY_EVENT_ID, 123);
```

- ✓ lto: { thin: true }
 - **LTO**: Link-Time Optimization it makes your final .so smaller and faster by optimizing across all source files during final linking.
 - thin: true: Use ThinLTO, a faster and more incremental version of LTO.

Simple explanation: Enables smarter compiler optimization *after* building the object files — but **thin** keeps build times reasonable.

Let me know if you want to see how libsurfaceflinger_sources is defined or what logtags look like.



The createSurfaceFlinger() mentioned above will call new SurfaceFlinger(), and then execute to: onFirstRef():

```
void SurfaceFlinger::onFirstRef()
{
   mEventQueue.init( this );
}
```

In onFirstRef(), the Handler is created and initialized: /frameworks/native/services/surfaceflinger/Scheduler/MessageQueue.cpp

Copy code
// MessageQueue.cpp
void MessageQueue::init(const sp<SurfaceFlinger>& flinger)



```
mFlinger = flinger;
  mLooper = new Looper( true );
  mHandler = new Handler(* this );
}
Copy code
    Then it will execute to SurfaceFlinger::init(), the main function of this method is:
Initializing EGL
Creating HWComposer
Initialize non-virtual display
Start the EventThread thread
Start boot animation
Copy code
// Do not call property_set on main thread which will be blocked by init
// Use StartPropertySetThread instead.
void SurfaceFlinger::init() {
  ALOGI( "SurfaceFlinger's main thread ready to run."
      "Initializing graphics H/W..." );
  Mutex::Autolock _l(mStateLock);
 // Set the CompositionEngine property and create a RenderEngine object
  // Get a RenderEngine for the given display / config (can't fail)
  // TODO(b/77156734): We need to stop casting and use HAL types when possible.
  // Sending maxFrameBufferAcquiredBuffers as the cache size is tightly tuned to single-
display.
  mCompositionEngine-> setRenderEngine(renderengine::RenderEngine::create(
      renderengine::RenderEngineCreationArgs::Builder()
         .setPixelFormat(static_cast <int32_t> (defaultCompositionPixelFormat))
         .setImageCacheSize(maxFrameBufferAcquiredBuffers)
         .setUseColorManagerment(useColorManagement)
        .setEnableProtectedContext(enable_protected_contents( false ))
         .setPrecacheToneMapperShaderOnly( false )
        .setSupportsBackgroundBlur(mSupportsBlur)
         .setContextPriority(useContextPriority
             ? renderengine::RenderEngine::ContextPriority::HIGH
             : renderengine::RenderEngine::ContextPriority::MEDIUM)
         .build()));
  mCompositionEngine -> setTimeStats(mTimeStats);
  LOG_ALWAYS_FATAL_IF(mVrFlingerRequestsDisplay,
      "Starting with vr flinger active is not currently supported." );
    //Create a HWComposer object and pass in a name attribute, then set the object properties
through mCompositionEngine->setHwComposer.
  mCompositionEngine->
```



```
setHwComposer(getFactory().createHWComposer(getBE().mHwcServiceName));
  mCompositionEngine->getHwComposer().setConfiguration( this ,
getBE().mComposerSequenceId);
    //processDisplayHotplugEventsLocked(); //Process any initial hotplug and resulting display
changes
    . //This method mainly calls initScheduler(displayId);
  // Process any initial hotplug and resulting display changes.
  processDisplayHotplugEventsLocked();
  const auto display = getDefaultDisplayDeviceLocked();
  LOG_ALWAYS_FATAL_IF(!display, "Missing internal display after registering composer
callback.");
  LOG_ALWAYS_FATAL_IF( !getHwComposer().isConnected(*display-> getId()),
             "Internal display is disconnected." );
  if (useVrFlinger) {
    auto vrFlingerRequestDisplayCallback = [ this ](bool requestDisplay) {
       // This callback is called from the vr flinger dispatch thread. We
       // need to call signalTransaction(), which requires holding
       // mStateLock when we're not on the main thread. Acquiring
       // mStateLock from the vr flinger dispatch thread trigger might a
       // deadlock in surface flinger (see b/66916578), so post a message
       // to be handled on the main thread instead.
      static cast< void >(schedule([= ] {
        ALOGI( "VR request display mode: requestDisplay=%d", requestDisplay);
        mVrFlingerRequestsDisplay = requestDisplay;
        signalTransaction();
      }));
    };
    mVrFlinger = dvr::VrFlinger::Create(getHwComposer(),getComposer(),
                        getHwComposer()
                            .fromPhysicalDisplayId( *display-> getId())
                            .value_or(0),
                        vrFlingerRequestDisplayCallback);
    if (! mVrFlinger) {
      ALOGE( "Failed to start vrflinger" );
    }
  }
  // initialize our drawing state
  mDrawingState = mCurrentState;
  // set initial conditions (eg unblank default device)
  initializeDisplays();
  char primeShaderCache[PROPERTY_VALUE_MAX];
```



```
property_get( "service.sf.prime_shader_cache", primeShaderCache, "1" );
  if (atoi(primeShaderCache)) {
    getRenderEngine().primeCache();
  }
  // Inform native graphics APIs whether the present timestamp is supported:
  const bool presentFenceReliable =
      ! getHwComposer().hasCapability(hal::Capability::PRESENT_FENCE_IS_NOT_RELIABLE);
  mStartPropertySetThread = getFactory().createStartPropertySetThread(presentFenceReliable);
  if (mStartPropertySetThread->Start() != NO_ERROR) {
    ALOGE( "Run StartPropertySetThread failed!" );
  }
  ALOGV( "Done initializing" );
}
Copy code
    First, let's see how to create
HWComposer:frameworks/native/services/surfaceflinger/SurfaceFlingerDefaultFactory.cpp
// make_unique is equivalent to new (can replace new and does not require a delete pointer,
which helps with code management).
std::unique_ptr<HWComposer> DefaultFactory::createHWComposer( const std::string&
serviceName) {
  return std::make_unique<android::impl::HWComposer> (serviceName);
}
    Initialize CompositionEngine:
std::unique ptr<compositionengine::CompositionEngine>
DefaultFactory::createCompositionEngine() {
   return compositionengine::impl::createCompositionEngine();
}
    The CompositionEngine object is also created through make_unique:
std::unique_ptr<compositionengine::CompositionEngine> createCompositionEngine() {
   return std::make_unique<CompositionEngine> ();
}
    Go back to flinger->init() processDisplayHotplugEventsLocked(); to handle the results of
any initial hot plug and display changes. In this method, the main call is
initScheduler(displayId):
Copy code
void SurfaceFlinger::initScheduler(DisplayId primaryDisplayId) {
  if (mScheduler) {
     // In practice it's not allowed to hotplug in/out the primary display once it's been
```



```
// connected during startup, but some tests do it, so just warn and return.
    ALOGW("Can't re-init scheduler");
    return;
  }
  auto currentConfig =
HwcConfigIndexType(getHwComposer().getActiveConfigIndex(primaryDisplayId));
  mRefreshRateConfigs =
      std::make_unique <scheduler::RefreshRateConfigs> (getHwComposer().getConfigs(
                                    primaryDisplayId),
                                currentConfig);
  mRefreshRateStats =
      std::make_unique <scheduler::RefreshRateStats>(*mRefreshRateConfigs, * mTimeStats,
                               currentConfig, hal::PowerMode::OFF);
  mRefreshRateStats -> setConfigMode(currentConfig);
  mPhaseConfiguration = getFactory().createPhaseConfiguration(* mRefreshRateConfigs);
    // Create a Scheduler object
  // start the EventThread
  mScheduler =
      getFactory().createScheduler([ this ](bool enabled) { setPrimaryVsyncEnabled(enabled); },
                      *mRefreshRateConfigs, * this );
  //Create app link and sf link
  mAppConnectionHandle =
      mScheduler ->createConnection("app", mPhaseConfiguration->
getCurrentOffsets().late.app,
                      impl::EventThread::InterceptVSyncsCallback());
  mSfConnectionHandle =
      mScheduler ->createConnection("sf", mPhaseConfiguration-> getCurrentOffsets().late.sf,
                      [ this ](nsecs_t timestamp) {
                        mInterceptor -> saveVSyncEvent(timestamp);
                     });
  mEventQueue ->setEventConnection(mScheduler->
getEventConnection(mSfConnectionHandle));
  mVSyncModulator.emplace( * mScheduler, mAppConnectionHandle, mSfConnectionHandle,
               mPhaseConfiguration -> getCurrentOffsets());
  mRegionSamplingThread =
       new RegionSamplingThread(* this, * mScheduler,
                   RegionSamplingThread::EnvironmentTimingTunables());
  // Dispatch a config change request for the primary display on scheduler
  // initialization, so that the EventThreads always contain a reference to a
  // prior configuration.
  //
```



```
// This is a bit hacky, but this avoids a back-pointer into the main SF
  // classes from EventThread, and there should be no run-time binder cost
  // anyway since there are no connected apps at this point.
  const nsecs_t vsyncPeriod =
      mRefreshRateConfigs -> getRefreshRateFromConfigId(currentConfig).getVsyncPeriod();
  mScheduler -> onPrimaryDisplayConfigChanged(mAppConnectionHandle,
primaryDisplayId.value,
                         currentConfig, vsyncPeriod);
}
Copy code
    See the links of app and sf in detail:
Copy code
Scheduler::ConnectionHandle Scheduler::createConnection(
     const_char * connectionName, nsecs_t phaseOffsetNs,
     impl::EventThread::InterceptVSyncsCallback interceptCallback) {
  auto vsyncSource = makePrimaryDispSyncSource(connectionName, phaseOffsetNs);
  auto eventThread = std::make_unique<impl::EventThread> (std::move(vsyncSource),
std::move(interceptCallback));
   return createConnection(std::move(eventThread));
Copy code
    You can see that the DispSyncSource object is created, and the construction method passes
in four values, dispSync object, phaseOffset offset, traceVsync is true, and name is app or sf
Copy code
 DispSyncSource::DispSyncSource(DispSync* dispSync, nsecs_t phaseOffset, bool traceVsync,
                  const char * name)
    : mName(name),
     mValue(base::StringPrintf("VSYNC-%s", name), 0), // mValue is assigned a value. The
VSYNC-app VSYNC-sf label we see on systrace is it
     mTraceVsync(traceVsync), // mTraceVsync is true, in the onDispSyncEvent method
     mVsyncOnLabel(base::StringPrintf( "VsyncOn-%s" , name)),
     mDispSync(dispSync),
     mPhaseOffset(base::StringPrintf( "VsyncOffset-%s", name), phaseOffset) //Initialize
mPhaseOffset When the vsync signal arrives, the offset of sf and app
Copy code
```

Great — you've traced a complete path from main() to SurfaceFlinger::init() and then deep into **HWComposer, RenderEngine, EventThread, and Scheduler setup**. Let's now break this entire process down into clear, simple explanations step-by-step. 🧠 🕏



🏁 1. SurfaceFlinger Starts Up



```
cpp
sp<SurfaceFlinger> flinger = surfaceflinger::createSurfaceFlinger();
```

- Calls new SurfaceFlinger()
- Because SurfaceFlinger inherits RefBase, its onFirstRef() gets called automatically.

$\ref{2.00}$ 2. onFirstRef() ightarrow Initialize Event Queue

```
void SurfaceFlinger::onFirstRef() {
    mEventQueue.init(this);
}
```

➤ What this does:

- Creates a Looper: for handling messages on the main thread.
- Creates a Handler: to post and process messages related to VSync or transactions.

3. SurfaceFlinger::init() — Main Initialization

This sets up all the **core graphics hardware/software components**.

🧱 Key Steps:

RenderEngine Creation

```
cpp
mCompositionEngine->setRenderEngine(RenderEngine::create(...));
```

- Talks directly to OpenGL / Vulkan drivers
- Used for rendering all frames and composition

✓ HWComposer Creation

```
mCompositionEngine->setHwComposer(getFactory().createHWComposer());
```

- Talks to hardware composer HAL (HWC) sends buffers for display
- Receives display events like hotplug

Handle Display Events

```
cpp
processDisplayHotplugEventsLocked();
```

- Detects connected displays (e.g., internal screen, HDMI)
- Initializes a default display device

4. Scheduler and VSync Configuration

```
initScheduler(displayId);
```

This sets up the **frame timing system**: when to draw, how often, how to sync with the display.

What Happens Inside:

➤ Create RefreshRateConfigs

```
cpp
mRefreshRateConfigs = new RefreshRateConfigs(...);
```

• Manages available display refresh rates (60Hz, 90Hz, etc.)

➤ Create Scheduler

```
mScheduler = getFactory().createScheduler(...);
```

• Controls frame pacing (when SF and apps draw)

• Uses VSync signals to maintain consistent timing

➤ Setup Connections

```
mAppConnectionHandle = mScheduler->createConnection("app", ...);
mSfConnectionHandle = mScheduler->createConnection("sf", ...);
```

- app: App layer (SurfaceView, OpenGL)
- **sf**: SurfaceFlinger internal work

Each one gets a **VSync source**, a **phase offset**, and an **event thread**.

➤ DispSyncSource

```
cpp
DispSyncSource::DispSyncSource(DispSync*, phaseOffset, true, name);
```

- Wraps a VSync signal with extra settings like when to fire callbacks
- The **offset** allows sf and app to draw at slightly different times (prevent jank)
- On **Systrace**, these appear as **VSYNC-app**, **VSYNC-sf** markers.

🧵 5. Threads and Message Loops

- EventThread: dispatches VSync to app or SF
- RegionSamplingThread: reads visible regions for blur/shaders
- StartPropertySetThread: sets display boot properties asynchronously

🚀 Final Boot Steps

- Draw first frame (boot animation)
- Prime shader cache
- Register Binder service: addService("SurfaceFlinger", ...)
- Start main run loop: flinger->run()

🔚 Summary Diagram (Simple View)

```
main()
createSurfaceFlinger()
onFirstRef()
Create Looper + Handler
init()
Create RenderEngine
Create HWComposer
Setup Displays
initScheduler()
Scheduler
VSync sources (app + sf)
EventThreads
Shader cache
StartPropertySetThread
```

Let me know if you want to walk through how the **app and SurfaceFlinger use VSync differently**, or what **EventThread** does in more detail.



```
* Draw first frame (boot animation) * Prime shader cache * Register Binder service:

`addService("SurfaceFlinger", ...)` * Start main run loop: `flinger->run()`

explain in details
```

Let's break down those final steps of SurfaceFlinger initialization in complete detail, covering:

- 1. **V** Drawing the first frame (boot animation)
- 2. 🚀 Priming the shader cache
- 3. Negistering the SurfaceFlinger Binder service
- 4. Starting the main run loop

1. Drawing the First Frame (Boot Animation)

Where it happens:

```
void SurfaceFlinger::init() {
    ...
    startBootAnim();
}
```

Why it's needed:



- When Android starts up, the boot animation (bootanimation.zip) is shown by an app called bootanimation.
- But before that, SurfaceFlinger needs to **draw something once**, or the screen might stay black.

How it works:

```
void SurfaceFlinger::startBootAnim() {
    property_set("service.bootanim.exit", "0");
    property_set("ctl.start", "bootanim");
}

// Let the bootanimation app run
// Start bootanimation service via init
}
```

- bootanimation is a small graphics app that uses SurfaceFlinger to draw to the screen.
- SurfaceFlinger must draw at least one black frame or bootanim may not render correctly.
- SurfaceFlinger posts this first frame using repaintEverything() or commitTransactions() internally, depending on mode.

🚀 2. Priming the Shader Cache

Where it happens:

```
void SurfaceFlinger::primeShaderCache() {
    ...
    renderengine->primeCache();
}
```

Why it's needed:

- When SurfaceFlinger uses **RenderEngine**, it uses **OpenGL/Vulkan shaders** for:
 - Blending
 - Blur
 - Rounded corners
 - Color correction, etc.
- Compiling shaders at runtime causes jank or stalls.
- So Android **pre-compiles and caches shaders** used by SurfaceFlinger.

How it works:

• RenderEngine has a **priming phase**:



```
void RenderEngine::primeCache() {
    // Draw dummy layers using all common shaders
}
```

- This ensures GPU driver compiles and links the shaders.
- Next time SF draws, everything is already in memory → no runtime delay.

🔪 3. Registering Binder Service

Where it happens:

```
SurfaceFlinger::init() {
    ...
    defaultServiceManager()->addService(String16("SurfaceFlinger"), this);
}
```

Why it's important:

- Apps and system services (like WindowManager) talk to SurfaceFlinger using Binder IPC.
- SurfaceFlinger needs to expose itself as a **system-wide Binder service**.

How it works:

- defaultServiceManager() gets the Binder system service manager.
- addService("SurfaceFlinger", this) registers SF with that name.
- Clients (like WindowManagerService) call:

```
java

IBinder flinger = ServiceManager.getService("SurfaceFlinger");
```

• Then they call methods like createSurfaceControl() or setTransactionState().

Without this step, no client could connect to SF — no windows, no rendering, nothing.

🔁 4. Start Main Run Loop

Where it happens:



```
cpp
flinger->run();
```

What it does:

- Enters a **forever loop**, waiting for events (like VSync, transaction commits, or display changes).
- This is powered by the **Looper** created in **onFirstRef()**.

• Internals:

```
void SurfaceFlinger::run() {
    do {
        mLooper->pollOnce(-1); // Block until there's something to do
    } while (!mShuttingDown);
}
```

➤ What kind of events does it wait for?

- **VSync callbacks** → do composition
- **Surface/Layer transactions** → apply new geometry/visibility
- **Hotplug events** → display connected/disconnected
- MessageQueue events → handle posted messages from apps or SF itself

This loop keeps SF alive and responsive **until the system shuts down**.

📕 Full Summary Table

| Step | Purpose | Code | Impact |
|--------------------------|--|--|---|
| ✓ Draw first frame | Avoid black screen, kick off boot animation | startBootAnim() | Boot animation visible via SurfaceFlinger |
| Prime shader cache | Precompile shaders to avoid jank | <pre>primeShaderCache()</pre> | Smooth future rendering |
| Negister Binder | Allow clients (e.g. WindowManager) to talk to SF | <pre>addService("SurfaceFlinger",)</pre> | Enables all window creation and composition logic |
| Start main run loop | Enter the poll loop for events and VSync | flinger->run() | Keeps SF alive and active for system lifecycle |



Would you like to now dive into what exactly happens **inside that main loop** when a frame is drawn? We can walk through **VSync** \rightarrow **surface draw** \rightarrow **composition** \rightarrow **HWComposer** \rightarrow **display** step-by-step next.



 \hookrightarrow

Prime shader what is this

The term **"Prime shader"** is not a standard or widely used concept in graphics programming (OpenGL, Vulkan, DirectX, etc.), Android, or shader languages (GLSL/HLSL). So, it could mean one of the following depending on **context**:

1. Custom Project Term or Nickname

- It might be a **project-specific term** used internally by a graphics engine, team, or OEM vendor.
- Example: Someone might call their "main rendering shader" or "important visual effect shader" a *Prime Shader* informally to imply it's the *most crucial* shader in the pipeline.

🔽 2. Misunderstanding or Typo

- It could be a typo or misinterpretation of:
 - "Primary Shader" (i.e., main fragment/vertex shader)
 - "Primitive Shader" (specific to AMD GPU architecture)

✓ 3. If Meant as "Primitive Shader" (AMD)

If you were reading AMD GPU documentation or graphics architecture notes, "Primitive Shader" is a real concept:

Primitive Shader (AMD GPUs)

- Used in AMD's GCN/RDNA architecture.
- A GPU stage that **combines vertex and geometry shader stages** for better performance.
- It processes:
 - Vertex shading
 - Culling



- Primitive assembly
- Purpose: **Improves performance** by reducing fixed-function stages.

But this is **only used in low-level GPU hardware** context, not Android APIs or OpenGL ES.

4. Main Shader in a Pipeline

- Sometimes "Prime shader" is used to informally refer to the **most important shader** in a graphics pipeline.
 - Example: In a deferred renderer, someone might refer to the G-buffer shader or final composite shader as the "prime shader."

Conclusion:

There's no official or standard "Prime Shader" in OpenGL, Vulkan, or Android graphics. If you saw this term in some code or documentation, please share the exact context — then I can give a precise answer.

Would you like to paste where you saw "Prime shader"? I'll explain what it means in that specific case.



 \hookrightarrow

`VSYNC-app`, `VSYNC-sf` what are these

The terms **VSYNC-app** and **VSYNC-sf** refer to **two different VSync timelines** used by the **Android graphics pipeline**.

These are part of how Android coordinates screen refreshes with app rendering and SurfaceFlinger (the system compositor). Let's break this down simply:

What is VSync?

VSync (Vertical Sync) is a signal from the display hardware that tells the system:

"I'm about to refresh the screen — if you have a new frame, now's the time to show it!"

In Android, the display usually runs at **60Hz or 120Hz**, so VSync fires every **16.6 ms or 8.3 ms**.





- This is the **VSync signal sent to the app** process.
- It tells your **application** (like a game or UI) when to start drawing its next frame.
- Used by Choreographer and RenderThread to kick off rendering.
- Think of it as:

"Hey App, prepare your next frame now so it's ready by the next screen refresh."

★ VSYNC-sf (SurfaceFlinger)

- This is the VSync signal sent to SurfaceFlinger, Android's system compositor.
- SurfaceFlinger collects all the app frames and composes them into a final image to send to the display.
- sf stands for SurfaceFlinger.
- Think of it as:

"Hey SurfaceFlinger, now it's your turn to compose the app frames and send the final image to the display."

Why Two Different VSyncs?

Because apps and SurfaceFlinger do different jobs and need different **timing**:

- The app gets an earlier signal so it has time to draw its frame.
- SurfaceFlinger gets a slightly later signal, **just in time to composite** all the app frames and push to the screen.

This **prevents frame drops**, avoids tearing, and keeps animation smooth.

Timeline (simplified view):

Where You See These?

You'll see **VSYNC-app** and **VSYNC-sf** in:



- dumpsys SurfaceFlinger output
- systrace / Perfetto timeline
- Android FrameTimeline logs

They help developers debug frame delays, jank, and latency.

Summary

| Term | Stands For | Who Uses It | Purpose |
|-----------|--------------------------|---------------------|----------------------------------|
| VSYNC-app | VSync for App | App / Choreographer | Tell app when to start rendering |
| VSYNC-sf | VSync for SurfaceFlinger | SurfaceFlinger | Tell SF when to start composing |

Would you like a visual timeline or real-world example using systrace?



Create an EventThread object and pass in the vsyncSource object associated with sf or app:

auto eventThread = std::make_unique<impl::EventThread>(std::move(vsyncSource),
std::move(interceptCallback));
illustrate:

- 1) Each ConnectionHandle object has an id, which is used as the key value of the mConnections property (map<id, Connection>) in the Scheduler object. The Connection object contains three properties: ConnectionHandle, EventThreadConnection, and EventThread.
- 2) In mScheduler->getEventConnection(mSfConnectionHandle), use the id of mSfConnectionHandle as the key value to find the corresponding Connection in Scheduler's mConnections (unordered_map<int64_t, Connection>) and return its EventThreadConnection member property.
- 3) In getHwComposer().registerCallback(), the registerCallback() methods of HwComposer and Device are called in sequence, and the SurfaceFlinger object is encapsulated into ComposerCallbackBridge in Device; for the encapsulated object, the registerCallback() methods of Composer and IComposerClient are called in sequence to inject it into the implementation class of IComposerClient.

Related questions:

① The screen refresh rate is faster than the system frame rate:



At this time, after the contents of the front buffer are fully mapped to the screen, the back buffer is not yet ready for the next frame. The screen will not be able to read the next frame, so it can only continue to display the graphics of the current frame, causing a frame to be displayed multiple times, which is called stuttering.

2 The system frame rate is faster than the screen refresh rate

. At this time, the screen has not completely mapped a frame from the front buffer to the screen, while the system has already prepared the next frame in the back buffer and requires the next frame to be read to the screen. This will cause the upper half of the screen to be the graphics of the previous frame, and the lower half to be the graphics of the next frame, causing multiple frames to be displayed on the screen, which is called screen tearing.

To solve the above problems, the Android display system generally has multiple levels of buffering, that is, while the screen is being refreshed, another buffer is used to prepare the next frame of data, thereby improving performance:

This explanation is about how Android handles **screen refresh synchronization** and how **SurfaceFlinger**, **EventThread**, and **VSync** coordination help manage smooth graphics rendering. Let's go through it piece by piece in **simple terms**.

What is EventThread?

EventThread is a class in Android's SurfaceFlinger that deals with vsync events.

- **VSync (Vertical Sync)** is a signal from the display hardware that tells the system: "Hey, I'm ready for the next frame."
- The EventThread receives these VSync signals and notifies apps or SurfaceFlinger to start rendering the next frame.
- The line:

```
cpp
auto eventThread = std::make_unique<impl::EventThread>(std::move(vsyncSource),
std::move(interceptCallback));
```

means: create a new EventThread object that listens to vsync signals from a vsyncSource.

How connections are managed

1. ConnectionHandle and Scheduler

• **ConnectionHandle** is like a **ticket or ID** that uniquely identifies a listener (like SurfaceFlinger or an app).



• These are stored in a map:

```
mConnections: map<id, Connection>
```

- Each Connection has:
 - ConnectionHandle the ID or key.
 - EventThreadConnection the link between client and EventThread.
 - EventThread the object that actually sends vsync signals.

2. How SurfaceFlinger gets a vsync connection

```
mScheduler->getEventConnection(mSfConnectionHandle)
```

- This asks the scheduler: "Give me the EventThreadConnection that belongs to SurfaceFlinger"
- It finds this by looking up the id in mConnections, and returns the EventThreadConnection.

3. Registering for VSync via HWC

```
cpp
getHwComposer().registerCallback()
```

- HwComposer (Hardware Composer HAL) is responsible for talking to the actual display hardware.
- When this method is called:
 - SurfaceFlinger wraps itself into a ComposerCallbackBridge.
 - This bridge is registered with the IComposerClient (which talks to HWC).
 - So now, when the screen's vsync event happens, HWC tells SurfaceFlinger through this callback.

Related Problems in Displaying Frames

① Screen refresh rate is higher than frame generation rate (i.e., screen is too fast)

- Display wants new frames quickly, but the app/system hasn't prepared them yet.
- Result: Same frame gets shown again → Stuttering.



② Frame generation rate is higher than screen refresh (i.e., system is too fast)

- App/system pushes new frame before screen is done with current one.
- Result: Screen shows part of two frames at once → Screen tearing (e.g., top half one frame, bottom half another).

How Android solves this: Buffering

- Uses multiple buffers:
 - Front buffer: What screen is displaying now.
 - Back buffer: What app is drawing next.
 - (Sometimes extra buffers too)
- While screen is using one buffer to display a frame, another buffer is used to **prepare the next frame**.
- This is called **double buffering** or **triple buffering**.
- Helps make the rendering **smooth** and avoids tearing or stuttering.

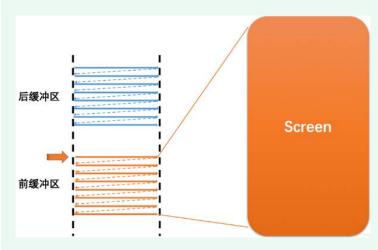
Summary

| Component | Purpose |
|--------------------|--|
| EventThread | Listens to VSync and sends events to apps/SurfaceFlinger |
| ConnectionHandle | Unique ID for a vsync listener |
| Scheduler | Manages all vsync connections |
| registerCallback() | Registers SurfaceFlinger to receive vsync from HWC |
| Multiple Buffers | Prevent screen tearing or stuttering |

Would you like a diagram for this whole flow or want me to break down a particular part in code?







Front buffer: frame buffer used to display content to the screen

Back buffer: frame buffer used to synthesize the next frame of graphics in the background

Vertical synchronization (VSync): A synchronization signal sent after the screen scans a frame from the buffer to the screen and before starting to scan the next frame. This signal is used to switch the front buffer and the back buffer.

Screen refresh rate (HZ): represents the number of times the screen refreshes in one second. Android phones generally refresh at 60HZ (that is, 60 frames per second, and 1 frame in about 16.67 milliseconds).

System frame rate (FPS): represents the number of frames synthesized by the system in one second. The size of this value is determined by the system algorithm and hardware.

This diagram illustrates the **double buffering** mechanism used in graphics rendering, especially relevant to **Android** and other modern systems. Let's break it down step by step with explanation of each part:

Image Explanation

On the **left** side of the image:

- The top half labeled "后缓冲区" (Back Buffer Region) in blue.
- The bottom half labeled "前缓冲区" (Front Buffer Region) in orange.
- Arrows show horizontal scanning of lines this simulates how a screen renders pixel rows.
- The **orange arrow** in the middle pointing right represents the **VSync signal** (Vertical Sync).
- The **right side**, a rectangular screen, shows that the **front buffer is currently being displayed**.

Terms and Concepts Explained



| Term | Meaning | |
|-------------------------------------|---|--|
| Front Buffer (前缓冲区) | This is the frame currently being shown on the screen. | |
| Back Buffer (后缓冲区) | This is the memory area where the next frame is being drawn in the background. It's not visible to the user yet. | |
| VSync (Vertical Synchronization) | A signal that tells the system "the screen has finished drawing the current frame." This is the right time to switch buffers (back \rightarrow front) to avoid visual glitches like tearing. | |
| Refresh Rate (Hz) | Number of times the screen is refreshed per second (e.g., 60Hz = 60 times per second). | |
| Frame Rate (FPS) | Number of frames the system is able to draw per second. This depends on CPU/GPU load and rendering complexity. | |

Why This Matters?

Without this system:

- You might see **tearing** part of a new frame gets shown before the old one is fully removed.
- Or **flickering**, if the frame is drawn while it's being shown.

So this buffering + VSync mechanism ensures:

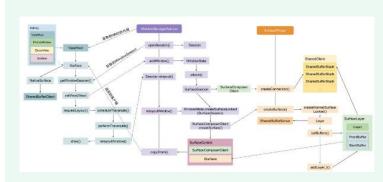
- Smooth animation
- Glitch-free rendering
- Synchronization between GPU-rendering and screen refresh

The Flow

- 1. The **GPU** draws the next frame into the **back buffer**.
- 2. The **screen** reads data from the **front buffer** line-by-line.
- 3. After the screen is done (VSync triggered), the buffers **swap**.
- 4. The newly drawn frame in the back buffer becomes the new front buffer.
- 5. The GPU starts working on the next frame in the now-empty back buffer.
- 6. Repeat this every 16.67ms (for 60Hz display).

If you'd like, I can also show how **triple buffering**, **adaptive refresh**, or **frame pacing** fit into this system.





4. Surface creation process

The diagram you've uploaded shows the **Surface creation flow** in Android — how an app creates a **Surface** to draw on, and how that **Surface** is eventually connected to **SurfaceFlinger**, Android's system compositor.

Let me break this down step by step, in simple and detailed terms:

1. App UI Setup

- An Activity creates UI components like ViewRoot, PhoneWindow, DecorView, etc.
- These are part of the Android view hierarchy.
- The bottom-most component here is the Surface.

2. Surface Creation Request (App Side)

- ViewRootImpl (shortened as ViewRoot) is the bridge between app UI and the system.
- It calls getWindowSession() to get a connection to the WindowManagerService (WMS).
- Then it uses **setView(View)** to attach the app's top-level view.

From there:

- requestLayout() → schedules a traversal.
- performTraversals() → performs layout, measure, and surface updates.
- It calls relayoutWindow() to tell WMS to create or resize a Surface.

🔃 3. WindowManagerService Side

- WindowManagerService handles windows and surfaces at the system level.
- openSession() → creates a Session object.
- addWindow() → adds the app's window to WMS.



• Session.relayout() → triggers relayoutWindow() on the server side.

Inside this, we go deeper:

- WindowState.createSurfaceLocked(SurfaceSession) is called.
- It uses **SurfaceComposerClient** to create a surface.

This is the point where Android starts talking to **SurfaceFlinger**.

4. SurfaceComposerClient to SurfaceFlinger

- SurfaceComposerClient.createSurface() connects to SurfaceFlinger and requests a new surface.
- Internally, SurfaceFlinger:
 - createConnection() → sets up a connection.
 - createSurface() → allocates a new surface.
 - A **SharedBufferServer** is created to manage buffer queues.

5. Shared Buffer Stack Setup

- SurfaceFlinger sets up:
 - SharedClient
 - SharedBufferStack
 - These stacks are buffer queues used for rendering.

A Layer is also created:

- createNormalSurfaceLocked() → allocates a Layer.
- setBuffers() → sets up front and back buffers.
- addLayer_l() → adds the layer to the SurfaceFlinger composition list.

6. Return Surface to App

- The app finally receives a **SurfaceControl** object.
- **SurfaceControl** wraps an **ISurface** interface and gives the app **low-level control** over the surface (position, visibility, buffer submission, etc.).

Now the app can:

• Use a Canvas or OpenGL to draw on the surface.



• The drawn frames are sent to SurfaceFlinger for final composition and display.

Summary (Simplified Flow):

Would you like a text version of the flow with just the important method calls, or a side-by-side comparison of roles between the App, WMS, and SurfaceFlinger?



confusing explain in simple words

Let's simplify this step-by-step.

Goal: How does an Android app get a Surface to draw its screen?

A **Surface** is like a **blank canvas** where the app draws its screen. That canvas is later shown on the display by **SurfaceFlinger**, Android's system compositor.

Step 1: App asks for a window

When your app starts:

- Android creates a window for it using the WindowManager.
- That window needs a surface (a place to draw).

Step 2: App asks for a Surface

The app (via ViewRootImpl) tells the system:

"Please give me a surface so I can draw my UI."

This is done using a system call named relayoutWindow().

Step 3: System (WindowManagerService) talks to SurfaceFlinger

The system's WindowManagerService:

• Calls SurfaceComposerClient, which talks to SurfaceFlinger.

Think of it like this:

"Hey SurfaceFlinger, I need a new surface for this app's window."

Step 4: SurfaceFlinger creates a Layer (like a sheet of glass)

SurfaceFlinger:

- Creates a **layer**, which is just a visual object to show on screen.
- Sets up **buffers** (memory areas to draw into).
- Gives back a **SurfaceControl**, which the app can use to control the surface.

🥙 Step 5: App uses the Surface to draw

Now the app:

- Draws on the surface using Canvas, OpenGL, or Vulkan.
- When it's done, the surface content is sent to **SurfaceFlinger**.

Step 6: SurfaceFlinger puts the surface on screen

SurfaceFlinger takes all surfaces (app, status bar, keyboard, etc.) and **composes** them into the final image.

It sends that image to the **display**.

Simple Analogy

| Component | Like in real life |
|----------------------|--|
| Арр | A painter |
| Surface | A canvas to draw on |
| WindowManagerService | The gallery manager who gives canvases to artists |
| SurfaceFlinger | The final framer who puts all canvases on the wall (display) |
| SurfaceControl | The handles on the canvas (lets you move or resize it) |

Let me know if you'd like a visual drawing to go with this.



The process of creating a Surface is the process of displaying the Activity. The specific implementation of Activity.makeVisible() is called in ActivityThread.handleResumeActivity():

```
Copy code
void makeVisible() {
  if (! mWindowAdded) {
    ViewManager wm = getWindowManager(); // Here getWindowManager gets the
WindowManagerImpl object
    wm.addView(mDecor, getWindow().getAttributes());
    mWindowAdded = true;
  }
  mDecor.setVisibility(View.VISIBLE);
}
Copy code
  WindowManagerImpl.java:
public void addView(@NonNull View view, @NonNull ViewGroup.LayoutParams params ) {
  applyDefaultToken( params );
  mGlobal.addView(view, params, mDisplay, mParentWindow);
}
  WindowManagerGlobal.java:
Copy code
public void addView(View view, ViewGroup.LayoutParams params, Display display, Window
parentWindow) {
  final WindowManager.LayoutParams wparams = (WindowManager.LayoutParams) params;
  // Create ViewRootImpl
  ViewRootImpl root = new ViewRootImpl(view.getContext(), display);
  view.setLayoutParams(wparams);
  mViews.add(view);
  mRoots.add(root);
  mParams.add(wparams);
  // Set View
  root.setView(view, wparams, panelParentView);
}
Copy code
  Create ViewRootImpl:
  public ViewRootImpl(Context context, Display display) {
```



```
// Get the proxy class of IWindowSession
    this (context, display, WindowManagerGlobal.getWindowSession(),
        false /* useSfChoreographer */ );
  WindowManagerGlobal.java:
Copy code
  @UnsupportedAppUsage
  public static IWindowSession getWindowSession() {
    synchronized (WindowManagerGlobal. class ) {
      if (sWindowSession == null ) {
        try {
          // Emulate the legacy behavior. The global instance of InputMethodManager
          // was instantiated here.
          // TODO(b/116157766): Remove this hack after cleaning up
@UnsupportedAppUsage
          // Get IMS The proxy class
          InputMethodManager.ensureDefaultInstanceForDefaultDisplayIfNecessary();
          IWindowManager windowManager = getWindowManagerService();
          // After calling Binder, finally call WMS
          sWindowSession = windowManager.openSession(
              new IWindowSessionCallback.Stub() {
                @Override
                public void onAnimatorScaleChanged( float scale) {
                 ValueAnimator.setDurationScale(scale);
               }
             });
       } catch (RemoteException e) {
          throw e.rethrowFromSystemServer();
       }
     }
     return sWindowSession;
   }
 }
Copy code
 WindowManagerService.openSession:
Copy code
  // -----
  // IWindowManager API
  // -----
  @Override
  public IWindowSession openSession(IWindowSessionCallback callback) {
    // Create session object
    return new Session(this, callback);
```

```
}
Copy code
  The data is written back to the app process through Binder again, and the obtained object is
the Session proxy object IWindowSession.
  After creating the ViewRootImpl object, call the setView() method of the object:
  ViewRootImpl:
Copy code
public void setView(View view, WindowManager.LayoutParams attrs, View panelParentView) {
 synchronized (this) {
  requestLayout(); // See the analysis below for details
  // Call through Binder and enter the Session of the system process
        res = mWindowSession.addToDisplayAsUser(mWindow, mSeq, mWindowAttributes,
                getHostVisibility(), mDisplay.getDisplayId(), userId, mTmpFrame,
                mAttachInfo.mContentInsets, mAttachInfo.mStableInsets,
                mAttachInfo.mDisplayCutout, inputChannel,
                mTempInsets, mTempControls);
}
Copy code
Session.java
Copy code
  @Override
  public int addToDisplayAsUser(IWindow window, int seq, WindowManager.LayoutParams
attrs,
       int viewVisibility, int displayId, int userId, Rect outFrame,
      Rect outContentInsets, Rect outStableInsets,
      DisplayCutout.ParcelableWrapper outDisplayCutout, InputChannel outInputChannel,
      InsetsState outInsetsState, InsetsSourceControl[] outActiveControls) {
    // Call the addWindow method of WMS
    return mService.addWindow( this , window, seg, attrs, viewVisibility, displayId, outFrame,
        outContentInsets, outStableInsets, outDisplayCutout, outInputChannel,
        outInsetsState, outActiveControls, userId);
  }
Copy code
WindowManagerService.java:
Copy code
  public int addWindow(Session session, IWindow client, int seq,
      LayoutParams attrs, int viewVisibility, int displayId, Rect outFrame,
```



```
Rect outContentInsets, Rect outStableInsets,
      DisplayCutout.ParcelableWrapper outDisplayCutout, InputChannel outInputChannel,
      InsetsState outInsetsState, InsetsSourceControl[] outActiveControls,
      int requestUserId) {
    Arrays.fill(outActiveControls, null);
    int [] appOp = new int [1];
    final boolean isRoundedCornerOverlay = (attrs.privateFlags
         & PRIVATE_FLAG_IS_ROUNDED_CORNERS_OVERLAY) != 0;
    int res = mPolicy.checkAddPermission(attrs.type, isRoundedCornerOverlay,
attrs.packageName,
        appOp);
    if (res != WindowManagerGlobal.ADD_OKAY) {
      return res;
   }
    WindowState parentWindow = null;
    final int callingUid = Binder.getCallingUid();
    final int callingPid = Binder.getCallingPid();
    final long origId = Binder.clearCallingIdentity();
    final int type = attrs.type;
    synchronized (mGlobalLock) {
      if (! mDisplayReady) {
         throw new IllegalStateException("Display has not been initialized");
      }
•••
      // Create WindowState
      final WindowState win = new WindowState(this, session, client, token, parentWindow,
          appOp[0], seq, attrs, viewVisibility, session.mUid, userId,
          session.mCanAddInternalSystemWindow);
      if (win.mDeathRecipient == null ) {
         // Client has apparently died, so there is no reason to
         // continue.
        ProtoLog.w(WM_ERROR, "Adding window client %s"
             + " that is dead, aborting.", client.asBinder());
         return WindowManagerGlobal.ADD APP EXITING;
      }
      if (win.getDisplayContent() == null ) {
        ProtoLog.w(WM ERROR, "Adding window to Display that has been removed.");
         return WindowManagerGlobal.ADD_INVALID_DISPLAY;
      }
      // Adjust WindowManager's LayoutParams parameters
      final DisplayPolicy displayPolicy = displayContent.getDisplayPolicy();
      displayPolicy.adjustWindowParamsLw(win, win.mAttrs, callingPid, callingUid);
```



```
res = displayPolicy.validateAddingWindowLw(attrs, callingPid, callingUid);
      if (res != WindowManagerGlobal.ADD_OKAY) {
         return res;
      }
      // Open the input channel
      final boolean openInputChannels = (outInputChannel!= null
          && (attrs.inputFeatures & INPUT_FEATURE_NO_INPUT_CHANNEL) == 0 );
      if (openInputChannels) {
        win.openInputChannel(outInputChannel);
      }
      displayContent.getInputMonitor().setUpdateInputWindowsNeededLw();
      boolean focusChanged = false;
      // When the window can receive key events, update the focused window
      if (win.canReceiveKeys()) {
        focusChanged = updateFocusedWindowLocked(UPDATE_FOCUS_WILL_ASSIGN_LAYERS,
             false /* updateInputWindows */);
        if (focusChanged) {
          imMayMove = false;
        }
      }
      if (imMayMove) {
        displayContent.computeImeTarget( true /* updateImeTarget */ );
      }
   }
    Binder.restoreCallingIdentity(origId);
    return res;
  }
Copy code
  Create a SurfaceSession object and add the current Session to the WMS.mSessions member
variable.
  Session.java:
Copy code
  void windowAddedLocked(String packageName) {
    mPackageName = packageName;
    mRelayoutTag = "relayoutWindow: " + mPackageName;
    if (mSurfaceSession == null) {
```



```
if (DEBUG) {
         Slog.v(TAG_WM, "First window added to " + this + ", creating SurfaceSession");
      }
      mSurfaceSession = new SurfaceSession();
      ProtoLog.i(WM_SHOW_TRANSACTIONS, " NEW SURFACE SESSION %s", mSurfaceSession);
      mService.mSessions.add(this);
       if (mLastReportedAnimatorScale != mService.getCurrentAnimatorScale()) {
         mService.dispatchNewAnimatorScaleLocked(this);
      }
    mNumWindow ++;
Copy code
  The creation of SurfaceSession will call JNI, and call nativeCreate() in JNI.
  android_view_SurfaceSession.cpp:
static ilong nativeCreate(JNIEnv* env, iclass clazz) {
  SurfaceComposerClient * client = new SurfaceComposerClient();
  client ->incStrong(( void * )nativeCreate);
  return reinterpret_cast<jlong> (client);
}
Copy code
static jlong nativeCreate(JNIEnv* env, jclass clazz, jobject sessionObj,
    jstring nameStr, jint w, jint h, jint format, jint flags, jlong parentObject,
    jobject metadataParcel) {
  ScopedUtfChars name(env, nameStr);
  sp <SurfaceComposerClient> client;
  if (sessionObj != NULL) {
    client = android_view_SurfaceSession_getClient(env, sessionObj);
  } else {
    client = SurfaceComposerClient::getDefault();
  }
  SurfaceControl *parent = reinterpret_cast<SurfaceControl*> (parentObject);
  sp <SurfaceControl> surface;
  LayerMetadata metadata;
  Parcel * parcel = parcelForJavaObject(env, metadataParcel);
  if (parcel && !parcel-> objectsCount()) {
    status_t err = metadata.readFromParcel(parcel);
     if (err != NO_ERROR) {
     jniThrowException(env, "java/lang/IllegalArgumentException",
                "Metadata parcel has wrong format" );
    }
  }
  status_t err = client-> createSurfaceChecked(
      String8(name.c_str()), w, h, format, & surface, flags, parent, std::move(metadata));
```

```
if (err == NAME NOT FOUND) {
    jniThrowException(env, "java/lang/IllegalArgumentException", NULL);
     return 0;
  } else if (err != NO_ERROR) {
    jniThrowException(env, OutOfResourcesException, NULL);
    return 0;
  }
  surface ->incStrong(( void * )nativeCreate);
  return reinterpret_cast<jlong> (surface.get());
}
Copy code
 Get the SurfaceComposerClient object through the above JNI interface as the proxy object for
communicating with SurfaceFlinger.
Copy code
void SurfaceComposerClient::onFirstRef() {
  // getComposerService() will return the BpSurfaceFlinger object
  sp<ISurfaceComposer> sf(ComposerService::getComposerService());
  if (sf != nullptr && mStatus == NO_INIT) {
    sp <ISurfaceComposerClient> conn;
    // Call SF's createConnection()
    conn = sf-> createConnection();
     if (conn != nullptr) {
      mClient = conn;
      mStatus = NO_ERROR;
    }
  }
}
Copy code
 For example, the screenshot interface will call the capture interface through SurfaceControl:
Copy code
status_t ScreenshotClient::capture( const sp<IBinder>& display, ui::Dataspace reqDataSpace,
                   ui::PixelFormat reqPixelFormat, const Rect& sourceCrop,
                   uint32_t reqWidth, uint32_t reqHeight, bool useIdentityTransform,
                   ui::Rotation rotation, bool captureSecureLayers,
                   sp <GraphicBuffer>* outBuffer, bool& outCapturedSecureLayers) {
  sp <ISurfaceComposer> s(ComposerService::getComposerService());
  if (s == nullptr) return NO_INIT;
  status_t ret = s-> captureScreen(display, outBuffer, outCapturedSecureLayers, reqDataSpace,
                    reqPixelFormat, sourceCrop, reqWidth, reqHeight,
                    useIdentityTransform, rotation, captureSecureLayers);
  if (ret != NO_ERROR) {
     return ret;
  }
```

```
return ret;
}
Copy code
Then take a closer look at the core SurfaceFlinger implementation:
SurfaceFlinger.cpp:
sp<ISurfaceComposerClient> SurfaceFlinger::createConnection() {
  // Create a Client
  const sp<Client> client = new Client( this );
  return client->initCheck() == NO_ERROR? client: nullptr;
}
  Back to the previous step, after creating the ViewRootImpl object, the setView() method of the
object is called. The requestLayout() method is called in setView(). Now let's take a closer look at
the method call process:
Copy code
  @Override
  public void requestLayout() {
     if (! mHandlingLayoutInLayoutRequest) {
      checkThread();
      mLayoutRequested = true;
      scheduleTraversals();
    }
  }
Copy code
Copy code
  @UnsupportedAppUsage
  void scheduleTraversals() {
     if (! mTraversalScheduled) {
      mTraversalScheduled = true;
      mTraversalBarrier = mHandler.getLooper().getQueue().postSyncBarrier();
              //Start TraversalRunnable
      mChoreographer.postCallback(
           Choreographer.CALLBACK_TRAVERSAL, mTraversalRunnable, null );
      notifyRendererOfFramePending();
      pokeDrawLockIfNeeded();
    }
  }
Copy code
  final class TraversalRunnable implements Runnable {
    @Override
    public void run() {
      doTraversal();
    }
  }
```



```
Copy code
  void doTraversal() {
     if (mTraversalScheduled) {
      mTraversalScheduled = false;
      mHandler.getLooper().getQueue().removeSyncBarrier(mTraversalBarrier);\\
      if (mProfile) {
         Debug.startMethodTracing( "ViewAncestor" );
      }
              //Call performTraversals
       performTraversals();
if (mProfile) {
         Debug.stopMethodTracing();
         mProfile = false;
      }
    }
  }
Copy code
Copy code
  private void performTraversals() {
    // cache mView since it is used so much below...
    final View host = mView; // It is DecorView
     if (mFirst || windowShouldResize || viewVisibilityChanged || cutoutChanged || params !=
null
         | | mForceNextWindowRelayout) {
      mForceNextWindowRelayout = false;
      if (isViewVisible) {
         // If this window is giving internal insets to the window
         // manager, and it is being added or changing its visibility,
         // then we want to first give the window manager "fake"
         // insets to cause it to effectively ignore the content of
         // the window during layout. This avoids it briefly causing
         // other windows to resize/move based on the raw frame of the
         // window, waiting until we can finish laying out this window
         // and get back to the window manager with the ultimately
         // computed insets.
         insetsPending = computesInternalInsets && (mFirst | | viewVisibilityChanged);
      }
      try {
         if (DEBUG_LAYOUT) {
           Log.i(mTag, "host=w:" + host.getMeasuredWidth() + ", h:" +
               host.getMeasuredHeight() + ", params=" + params);
         }
```



```
if (mAttachInfo.mThreadedRenderer!= null) {
           // relayoutWindow may decide to destroy mSurface. As that decision
           // happens in WindowManager service, we need to be defensive here
           // and stop using the surface in case it gets destroyed.
          if (mAttachInfo.mThreadedRenderer.pause()) {
             // Animations were running so we need to push a frame
             // to resume them
            mDirty.set(0, 0, mWidth, mHeight);
          }
mChoreographer.mFrameInfo.addFlags(FrameInfo.FLAG_WINDOW_LAYOUT_CHANGED);
        // Key function relayoutWindow
        relayoutResult = relayoutWindow(params, viewVisibility, insetsPending);
        if (DEBUG_LAYOUT) Log.v(mTag, "relayout: frame=" + frame.toShortString()
             + " cutout=" + mPendingDisplayCutout.get().toString()
             + " surface=" + mSurface);
        // If the pending {@link MergedConfiguration} handed back from
         // {@link #relayoutWindow} does not match the one last reported,
         // WindowManagerService has reported back a frame from a configuration not yet
         // handled by the client. In this case, we need to accept the configuration so we
         // do not lay out and draw with the wrong configuration.
        if (! mPendingMergedConfiguration.equals(mLastReportedMergedConfiguration)) {
           if (DEBUG_CONFIGURATION) Log.v(mTag, "Visible with new config: "
               + mPendingMergedConfiguration.getMergedConfiguration());
          performConfigurationChange(mPendingMergedConfiguration, ! mFirst,
               INVALID_DISPLAY /* same display */);
          updatedConfiguration = true;
        }
      } catch (RemoteException e) {
      }
    }
    boolean cancelDraw = mAttachInfo.mTreeObserver.dispatchOnPreDraw() | | ! isViewVisible;
    if (! cancelDraw) {
       if (mPendingTransitions != null && mPendingTransitions.size() > 0 ) {
         for (int i = 0; i < mPendingTransitions.size(); ++ i) {
          mPendingTransitions.get(i).startChangingAnimations();
        mPendingTransitions.clear();
```

```
// Start drawing, which calls draw(fullRedrawNeeded);
      performDraw();
    } else {
       if (isViewVisible) {
         // Try again
         scheduleTraversals();
      } else if (mPendingTransitions != null && mPendingTransitions.size() > 0 ) {
         for (int i = 0; i < mPendingTransitions.size(); ++ i) {
           mPendingTransitions.get(i).endChangingAnimations();
         mPendingTransitions.clear();
      }
    }
    if (mAttachInfo.mContentCaptureEvents != null ) {
      notifyContentCatpureEvents();
    }
    mIsInTraversal = false;
  }
Copy code
  Let's look at the implementation of relayoutWindow:
Copy code
    if (mSurfaceControl.isValid()) {
       if (! useBLAST()) {
         // First create a local Surface, then call copyFrom to copy the SurfaceControl
information to Surface
         mSurface.copyFrom(mSurfaceControl);
      } else {
         final Surface blastSurface = getOrCreateBLASTSurface(mSurfaceSize.x,
           mSurfaceSize.y);
         // If blastSurface == null that means it hasn't changed since the last time we
         // called. In this situation, avoid calling transferFrom as we would then
         // inc the generation ID and cause EGL resources to be recreated.
         if (blastSurface != null ) {
           mSurface.transferFrom(blastSurface);
         }
      }
    } else {
      destroySurface();
    }
```



Copy code

The SurfaceControl class can be considered as a wrapper class, which will eventually execute copyFrom() to return it to the App client:

```
Copy code
  @UnsupportedAppUsage
  public void copyFrom(SurfaceControl other) {
    if (other == null) {
       throw new IllegalArgumentException("other must not be null" );
    }
    long surfaceControlPtr = other.mNativeObject;
    if (surfaceControlPtr == 0) {
       throw new NullPointerException(
           "null SurfaceControl native object. Are you using a released SurfaceControl?" );
    }
    // Get the source SurfaceControl through JNI
    long newNativeObject = nativeGetFromSurfaceControl(mNativeObject, surfaceControlPtr);
    synchronized (mLock) {
       if (newNativeObject == mNativeObject) {
         return;
      }
      if (mNativeObject != 0) {
        nativeRelease(mNativeObject);
      }
      // Save to global mNativeObject for external call
      setNativeObjectLocked(newNativeObject);
  }
Copy code
```

Thanks for sharing the image. Let's now go over it in **simple terms**, with each part explained like a story:

What you're looking at:

This image shows how **SurfaceFlinger** works inside Android to draw everything on the screen — from apps, status bar, keyboard, etc.

Key Layers in the Diagram:

1. App Process (left box)



- This is your Android app (e.g., Facebook, YouTube).
- It wants to show something on the screen.

It uses:

- Surface to draw on.
- SurfaceControl to talk to SurfaceFlinger (to say things like move, hide, resize).

2. SurfaceFlinger Process (middle box)

This is a **system service** whose job is to **collect surfaces from all apps** and **compose** (combine) them into one image to show on the screen.

K It handles:

- Layers: One for each app or UI piece (like status bar).
- Composition: Blending all layers into a final image.
- Hardware Composer (HWC): Sends the final image to display.

3. Hardware Composer (right box)

- This is either hardware or driver software.
- Gets the final image and shows it on your display.

o Main Flow (Simple Version):

1. App creates a surface

- App says: "I need a canvas to draw on."
- Gets a Surface and a SurfaceControl.

2. App draws something

- Maybe some buttons, images, text.
- Draws it using OpenGL, Vulkan, Canvas, etc.

3. App sends the drawing to SurfaceFlinger

• When drawing is done, it's passed to SurfaceFlinger.

4. SurfaceFlinger collects all surfaces

- From your app
- From system UI (status bar, keyboard, etc.)
- From animations



5. SurfaceFlinger composes them together

- Like layers in Photoshop, it puts them in order.
- It clips, scales, blends them.

6. Sends final image to Hardware Composer

• The hardware shows the final frame on your screen.

Why SurfaceControl?

- Think of SurfaceControl as a remote control to manage the surface:
 - Move it
 - Resize it
 - Change transparency
 - Stack above/below others

The app uses this to **tell SurfaceFlinger how to treat the surface**, but the app **can't draw directly with it** — drawing is done via **Surface**.

🖈 Summary (like a team):

| Component | Role | |
|-------------------|---|--|
| Арр | Creates surface & draws content | |
| Surface | Canvas (where drawing happens) | |
| SurfaceControl | Remote control to tell SurfaceFlinger how to place/show the surface | |
| SurfaceFlinger | The painter-in-chief: collects surfaces and makes a final image | |
| Hardware Composer | The delivery guy: sends final image to screen | |

Let me know if you'd like a **simplified sketch** of this flow.



Surfaceflinger-2



The process of creating a Surface is the process of displaying the Activity. The specific implementation of Activity.makeVisible() is called in ActivityThread.handleResumeActivity():

```
Copy code
void makeVisible() {
  if (! mWindowAdded) {
    ViewManager wm = getWindowManager(); // Here getWindowManager gets the
WindowManagerImpl object
    wm.addView(mDecor, getWindow().getAttributes());
    mWindowAdded = true;
  }
  mDecor.setVisibility(View.VISIBLE);
}
Copy code
  WindowManagerImpl.java:
public void addView(@NonNull View view, @NonNull ViewGroup.LayoutParams params ) {
  applyDefaultToken( params );
  mGlobal.addView(view, params, mDisplay, mParentWindow);
}
  WindowManagerGlobal.java:
Copy code
public void addView(View view, ViewGroup.LayoutParams params, Display display, Window
parentWindow) {
  final WindowManager.LayoutParams wparams = (WindowManager.LayoutParams) params;
  // Create ViewRootImpl
  ViewRootImpl root = new ViewRootImpl(view.getContext(), display);
  view.setLayoutParams(wparams);
  mViews.add(view);
  mRoots.add(root);
  mParams.add(wparams);
  // Set View
  root.setView(view, wparams, panelParentView);
}
Copy code
  Create ViewRootImpl:
```



```
public ViewRootImpl(Context context, Display display) {
    // Get the proxy class of IWindowSession
    this (context, display, WindowManagerGlobal.getWindowSession(),
         false /* useSfChoreographer */);
  }
  WindowManagerGlobal.java:
Copy code
  @UnsupportedAppUsage
  public static IWindowSession getWindowSession() {
    synchronized (WindowManagerGlobal. class ) {
      if (sWindowSession == null ) {
        try {
           // Emulate the legacy behavior. The global instance of InputMethodManager
           // was instantiated here.
           // TODO(b/116157766): Remove this hack after cleaning up
@UnsupportedAppUsage
           // Get IMS The proxy class
          InputMethodManager.ensureDefaultInstanceForDefaultDisplayIfNecessary();
          IWindowManager windowManager = getWindowManagerService();
           // After calling Binder, finally call WMS
          sWindowSession = windowManager.openSession(
               new IWindowSessionCallback.Stub() {
                 @Override
                 public void onAnimatorScaleChanged( float scale) {
                   ValueAnimator.setDurationScale(scale);
                }
              });
        } catch (RemoteException e) {
           throw e.rethrowFromSystemServer();
        }
      return sWindowSession;
    }
  }
Copy code
 WindowManagerService.openSession:
Copy code
  // IWindowManager API
  @Override
  public IWindowSession openSession(IWindowSessionCallback callback) {
```



```
// Create session object
    return new Session(this, callback);
  }
Copy code
  The data is written back to the app process through Binder again, and the obtained object is
the Session proxy object IWindowSession.
  After creating the ViewRootImpl object, call the setView() method of the object:
  ViewRootImpl:
Copy code
public void setView(View view, WindowManager.LayoutParams attrs, View panelParentView) {
 synchronized (this) {
  requestLayout(); // See the analysis below for details
  // Call through Binder and enter the Session of the system process
        res = mWindowSession.addToDisplayAsUser(mWindow, mSeq, mWindowAttributes,
                getHostVisibility(), mDisplay.getDisplayId(), userId, mTmpFrame,
                mAttachInfo.mContentInsets, mAttachInfo.mStableInsets,
                mAttachInfo.mDisplayCutout, inputChannel,
                mTempInsets, mTempControls);
 }
Copy code
Session.java
Copy code
  @Override
  public int addToDisplayAsUser(IWindow window, int seq, WindowManager.LayoutParams
attrs,
       int viewVisibility, int displayId, int userId, Rect outFrame,
      Rect outContentInsets, Rect outStableInsets,
      DisplayCutout.ParcelableWrapper outDisplayCutout, InputChannel outInputChannel,
      InsetsState outInsetsState, InsetsSourceControl[] outActiveControls) {
    // Call the addWindow method of WMS
    return mService.addWindow( this , window, seq, attrs, viewVisibility, displayId, outFrame,
        outContentInsets, outStableInsets, outDisplayCutout, outInputChannel,
        outInsetsState, outActiveControls, userId);
  }
Copy code
WindowManagerService.java:
Copy code
```



```
public int addWindow(Session session, IWindow client, int seg,
      LayoutParams attrs, int viewVisibility, int displayId, Rect outFrame,
      Rect outContentInsets, Rect outStableInsets,
      DisplayCutout.ParcelableWrapper outDisplayCutout, InputChannel outInputChannel,
      InsetsState outInsetsState, InsetsSourceControl[] outActiveControls,
      int requestUserId) {
    Arrays.fill(outActiveControls, null);
    int [] appOp = new int [1];
    final boolean isRoundedCornerOverlay = (attrs.privateFlags
         & PRIVATE FLAG IS ROUNDED CORNERS OVERLAY) != 0;
    int res = mPolicy.checkAddPermission(attrs.type, isRoundedCornerOverlay,
attrs.packageName,
        appOp);
    if (res != WindowManagerGlobal.ADD OKAY) {
       return res;
    }
    WindowState parentWindow = null;
    final int callingUid = Binder.getCallingUid();
    final int callingPid = Binder.getCallingPid();
    final long origId = Binder.clearCallingIdentity();
    final int type = attrs.type;
    synchronized (mGlobalLock) {
      if (! mDisplayReady) {
         throw new IllegalStateException("Display has not been initialized");
      }
      // Create WindowState
      final WindowState win = new WindowState( this , session, client, token, parentWindow,
          appOp[0], seq, attrs, viewVisibility, session.mUid, userId,
          session.mCanAddInternalSystemWindow);
      if (win.mDeathRecipient == null ) {
         // Client has apparently died, so there is no reason to
         // continue.
        ProtoLog.w(WM ERROR, "Adding window client %s"
             + " that is dead, aborting.", client.asBinder());
         return WindowManagerGlobal.ADD_APP_EXITING;
      }
      if (win.getDisplayContent() == null ) {
        ProtoLog.w(WM_ERROR, "Adding window to Display that has been removed." );
         return WindowManagerGlobal.ADD_INVALID_DISPLAY;
      }
      // Adjust WindowManager's LayoutParams parameters
```



```
final DisplayPolicy displayPolicy = displayContent.getDisplayPolicy();
      displayPolicy.adjustWindowParamsLw(win, win.mAttrs, callingPid, callingUid);
      res = displayPolicy.validateAddingWindowLw(attrs, callingPid, callingUid);
      if (res != WindowManagerGlobal.ADD_OKAY) {
         return res;
      }
      // Open the input channel
      final boolean openInputChannels = (outInputChannel!= null
          && (attrs.inputFeatures & INPUT_FEATURE_NO_INPUT_CHANNEL) == 0 );
      if (openInputChannels) {
        win.openInputChannel(outInputChannel);
      }
      displayContent.getInputMonitor().setUpdateInputWindowsNeededLw();
      boolean focusChanged = false;
      // When the window can receive key events, update the focused window
      if (win.canReceiveKeys()) {
        focusChanged = updateFocusedWindowLocked(UPDATE_FOCUS_WILL_ASSIGN_LAYERS,
             false /* updateInputWindows */);
        if (focusChanged) {
          imMayMove = false;
        }
      }
      if (imMayMove) {
        displayContent.computeImeTarget( true /* updateImeTarget */ );
      }
    }
    Binder.restoreCallingIdentity(origId);
    return res;
  }
Copy code
  Create a SurfaceSession object and add the current Session to the WMS.mSessions member
variable.
  Session.java:
Copy code
  void windowAddedLocked(String packageName) {
    mPackageName = packageName;
```



```
mRelayoutTag = "relayoutWindow: " + mPackageName;
     if (mSurfaceSession == null ) {
       if (DEBUG) {
         Slog.v(TAG_WM, "First window added to " + this + ", creating SurfaceSession" );
      }
      mSurfaceSession = new SurfaceSession();
      ProtoLog.i(WM_SHOW_TRANSACTIONS, " NEW SURFACE SESSION %s", mSurfaceSession);
      mService.mSessions.add(this);
       if (mLastReportedAnimatorScale != mService.getCurrentAnimatorScale()) {
         mService.dispatchNewAnimatorScaleLocked(this);
      }
    mNumWindow ++;
  }
Copy code
  The creation of SurfaceSession will call JNI, and call nativeCreate() in JNI.
  android_view_SurfaceSession.cpp:
static jlong nativeCreate(JNIEnv* env, jclass clazz) {
  SurfaceComposerClient * client = new SurfaceComposerClient();
  client ->incStrong(( void * )nativeCreate);
  return reinterpret_cast<jlong> (client);
}
Copy code
static ilong nativeCreate(JNIEnv* env, iclass clazz, jobject sessionObj,
    jstring nameStr, jint w, jint h, jint format, jint flags, jlong parentObject,
    jobject metadataParcel) {
  ScopedUtfChars name(env, nameStr);
  sp <SurfaceComposerClient> client;
  if (sessionObj != NULL) {
    client = android_view_SurfaceSession_getClient(env, sessionObj);
  } else {
    client = SurfaceComposerClient::getDefault();
  SurfaceControl *parent = reinterpret_cast<SurfaceControl*> (parentObject);
  sp <SurfaceControl> surface;
  LayerMetadata metadata;
  Parcel * parcel = parcelForJavaObject(env, metadataParcel);
  if (parcel && !parcel-> objectsCount()) {
    status_t err = metadata.readFromParcel(parcel);
     if (err != NO_ERROR) {
     iniThrowException(env, "java/lang/IllegalArgumentException",
                "Metadata parcel has wrong format" );
    }
  }
```

```
status t err = client-> createSurfaceChecked(
      String8(name.c_str()), w, h, format, & surface, flags, parent, std::move(metadata));
  if (err == NAME_NOT_FOUND) {
    jniThrowException(env, "java/lang/IllegalArgumentException", NULL);
     return 0;
  } else if (err != NO_ERROR) {
    jniThrowException(env, OutOfResourcesException, NULL);
    return 0;
  }
  surface ->incStrong(( void * )nativeCreate);
  return reinterpret_cast<jlong> (surface.get());
}
Copy code
 Get the SurfaceComposerClient object through the above JNI interface as the proxy object for
communicating with SurfaceFlinger.
Copy code
void SurfaceComposerClient::onFirstRef() {
  // getComposerService() will return the BpSurfaceFlinger object
  sp<ISurfaceComposer> sf(ComposerService::getComposerService());
  if (sf != nullptr && mStatus == NO_INIT) {
    sp <ISurfaceComposerClient> conn;
    // Call SF's createConnection()
    conn = sf-> createConnection();
     if (conn != nullptr) {
      mClient = conn;
      mStatus = NO_ERROR;
    }
  }
}
Copy code
 For example, the screenshot interface will call the capture interface through SurfaceControl:
Copy code
status t ScreenshotClient::capture( const sp<IBinder>& display, ui::Dataspace regDataSpace,
                   ui::PixelFormat reqPixelFormat, const Rect& sourceCrop,
                   uint32_t reqWidth, uint32_t reqHeight, bool useIdentityTransform,
                   ui::Rotation rotation, bool captureSecureLayers,
                   sp <GraphicBuffer>* outBuffer, bool& outCapturedSecureLayers) {
  sp <ISurfaceComposer> s(ComposerService::getComposerService());
  if (s == nullptr) return NO_INIT;
  status_t ret = s-> captureScreen(display, outBuffer, outCapturedSecureLayers, reqDataSpace,
                    regPixelFormat, sourceCrop, regWidth, regHeight,
                    useIdentityTransform, rotation, captureSecureLayers);
  if (ret != NO_ERROR) {
```

```
return ret;
  }
  return ret;
}
Copy code
Then take a closer look at the core SurfaceFlinger implementation:
SurfaceFlinger.cpp:
sp<ISurfaceComposerClient> SurfaceFlinger::createConnection() {
  // Create a Client
  const sp<Client> client = new Client( this );
  return client->initCheck() == NO_ERROR ? client : nullptr;
}
  Back to the previous step, after creating the ViewRootImpl object, the setView() method of the
object is called. The requestLayout() method is called in setView(). Now let's take a closer look at
the method call process:
Copy code
  @Override
  public void requestLayout() {
     if (! mHandlingLayoutInLayoutRequest) {
      checkThread();
      mLayoutRequested = true;
      scheduleTraversals();
    }
  }
Copy code
Copy code
  @UnsupportedAppUsage
  void scheduleTraversals() {
     if (! mTraversalScheduled) {
      mTraversalScheduled = true;
      mTraversalBarrier = mHandler.getLooper().getQueue().postSyncBarrier();
              //Start TraversalRunnable
      mChoreographer.postCallback(
           Choreographer.CALLBACK_TRAVERSAL, mTraversalRunnable, null );
      notifyRendererOfFramePending();
      pokeDrawLockIfNeeded();
    }
  }
Copy code
  final class TraversalRunnable implements Runnable {
    @Override
    public void run() {
      doTraversal();
```



```
}
Copy code
  void doTraversal() {
     if (mTraversalScheduled) {
      mTraversalScheduled = false;
      mHandler.getLooper().getQueue().removeSyncBarrier(mTraversalBarrier);
      if (mProfile) {
         Debug.startMethodTracing( "ViewAncestor" );
      }
              //Call performTraversals
       performTraversals();
if (mProfile) {
         Debug.stopMethodTracing();
         mProfile = false;
      }
    }
  }
Copy code
Copy code
  private void performTraversals() {
     // cache mView since it is used so much below...
    final View host = mView: // It is DecorView
     if (mFirst || windowShouldResize || viewVisibilityChanged || cutoutChanged || params !=
null
         | | mForceNextWindowRelayout) {
      mForceNextWindowRelayout = false;
      if (isViewVisible) {
         // If this window is giving internal insets to the window
         // manager, and it is being added or changing its visibility,
         // then we want to first give the window manager "fake"
         // insets to cause it to effectively ignore the content of
         // the window during layout. This avoids it briefly causing
         // other windows to resize/move based on the raw frame of the
         // window, waiting until we can finish laying out this window
         // and get back to the window manager with the ultimately
         // computed insets.
         insetsPending = computesInternalInsets && (mFirst |  viewVisibilityChanged);
      }
...
      try {
         if (DEBUG_LAYOUT) {
           Log.i(mTag, "host=w:" + host.getMeasuredWidth() + ", h:" +
```



```
host.getMeasuredHeight() + ", params=" + params);
        }
        if (mAttachInfo.mThreadedRenderer != null ) {
           // relayoutWindow may decide to destroy mSurface. As that decision
           // happens in WindowManager service, we need to be defensive here
           // and stop using the surface in case it gets destroyed.
          if (mAttachInfo.mThreadedRenderer.pause()) {
             // Animations were running so we need to push a frame
             // to resume them
            mDirty.set(0, 0, mWidth, mHeight);
          }
mChoreographer.mFrameInfo.addFlags(FrameInfo.FLAG_WINDOW_LAYOUT_CHANGED);
        // Key function relayoutWindow
        relayoutResult = relayoutWindow(params, viewVisibility, insetsPending);
        if (DEBUG_LAYOUT) Log.v(mTag, "relayout: frame=" + frame.toShortString()
             + " cutout=" + mPendingDisplayCutout.get().toString()
             + " surface=" + mSurface);
        // If the pending {@link MergedConfiguration} handed back from
         // {@link #relayoutWindow} does not match the one last reported,
         // WindowManagerService has reported back a frame from a configuration not yet
         // handled by the client. In this case, we need to accept the configuration so we
         // do not lay out and draw with the wrong configuration.
        if (! mPendingMergedConfiguration.equals(mLastReportedMergedConfiguration)) {
           if (DEBUG_CONFIGURATION) Log.v(mTag, "Visible with new config: "
               + mPendingMergedConfiguration.getMergedConfiguration());
          performConfigurationChange(mPendingMergedConfiguration, ! mFirst,
               INVALID_DISPLAY /* same display */ );
          updatedConfiguration = true;
        }
      } catch (RemoteException e) {
    }
    boolean cancelDraw = mAttachInfo.mTreeObserver.dispatchOnPreDraw() | | ! isViewVisible;
    if (! cancelDraw) {
       if (mPendingTransitions != null && mPendingTransitions.size() > 0 ) {
         for (int i = 0; i < mPendingTransitions.size(); ++ i) {
          mPendingTransitions.get(i).startChangingAnimations();
```



```
mPendingTransitions.clear();
      }
      // Start drawing, which calls draw(fullRedrawNeeded);
      performDraw();
    } else {
       if (isViewVisible) {
         // Try again
         scheduleTraversals();
      } else if (mPendingTransitions != null && mPendingTransitions.size() > 0 ) {
         for (int i = 0; i < mPendingTransitions.size(); ++ i) {
           mPendingTransitions.get(i).endChangingAnimations();
         }
         mPendingTransitions.clear();
      }
    }
    if (mAttachInfo.mContentCaptureEvents != null ) {
      notifyContentCatpureEvents();
    }
    mIsInTraversal = false;
  }
Copy code
  Let's look at the implementation of relayoutWindow:
Copy code
    if (mSurfaceControl.isValid()) {
       if (! useBLAST()) {
         // First create a local Surface, then call copyFrom to copy the SurfaceControl
information to Surface
         mSurface.copyFrom(mSurfaceControl);
      } else {
         final Surface blastSurface = getOrCreateBLASTSurface(mSurfaceSize.x,
           mSurfaceSize.y);
         // If blastSurface == null that means it hasn't changed since the last time we
         // called. In this situation, avoid calling transferFrom as we would then
         // inc the generation ID and cause EGL resources to be recreated.
         if (blastSurface != null ) {
           mSurface.transferFrom(blastSurface);
         }
      }
    } else {
      destroySurface();
```



```
Copy code
  The SurfaceControl class can be considered as a wrapper class, which will eventually execute
copyFrom() to return it to the App client:
Copy code
  @UnsupportedAppUsage
  public void copyFrom(SurfaceControl other) {
    if (other == null) {
       throw new IllegalArgumentException("other must not be null");
    }
    long surfaceControlPtr = other.mNativeObject;
    if (surfaceControlPtr == 0) {
       throw new NullPointerException(
           "null SurfaceControl native object. Are you using a released SurfaceControl?" );
    }
    // Get the source SurfaceControl through JNI
    long newNativeObject = nativeGetFromSurfaceControl(mNativeObject, surfaceControlPtr);
    synchronized (mLock) {
       if (newNativeObject == mNativeObject) {
         return;
      }
      if (mNativeObject != 0 ) {
        nativeRelease(mNativeObject);
      // Save to global mNativeObject for external call
      setNativeObjectLocked(newNativeObject);
    }
  }
Copy code
```

You've traced the entire **Surface creation pipeline** in Android—from the app process to **WindowManagerService** and back—beautifully. Let's break this down and explain **each key step** in **simple words**, with a clear focus on the flow of how a **Surface** is created and shown when an **Activity** becomes visible.

Summary: What is Happening?

When your Activity is resumed, Android goes through a complex pipeline that ends in creating a **Surface** using **SurfaceFlinger**, assigning it to your window, and letting it draw. Here's a simple picture:



```
nginx

Activity → WindowManager → ViewRootImpl → SurfaceSession → SurfaceControl → SurfaceFlinger → Surface
```

Step-by-Step Breakdown in Simple Terms

1. Activity Resumes → Calls makeVisible()

```
if (!mWindowAdded) {
    wm.addView(mDecor, getWindow().getAttributes());
}
```

- The **DecorView** (top-level View of your Activity) is added to the **WindowManager**.
- This triggers the whole Surface creation pipeline.

2. WindowManagerGlobal.addView()

```
java

ViewRootImpl root = new ViewRootImpl(view.getContext(), display);
root.setView(view, wparams, panelParentView);
```

- ViewRootImpl is created. It's a bridge between your app's UI and the windowing system.
- It handles **layout**, **drawing**, **input**, and **Surface management**.

3. ViewRootImpl.setView() calls mWindowSession.addToDisplayAsUser()

- mWindowSession is a proxy interface (IWindowSession) which talks to the WindowManagerService (WMS) via Binder.
- This tells the system: "I want to add a window and need a surface."

4. WMS receives request: addWindow()

```
java

final WindowState win = new WindowState(...);
win.openInputChannel(outInputChannel);
```

• WMS creates a **WindowState**, representing your app's window.

- It checks permissions, sets up focus, and prepares input.
- Now it needs a Surface.

5. Session's windowAddedLocked() → SurfaceSession is created

```
java
mSurfaceSession = new SurfaceSession();
```

- SurfaceSession is a low-level object that manages SurfaceControl objects.
- This step goes into JNI.

6. JNI: nativeCreate() → **creates** a SurfaceComposerClient

```
cpp
SurfaceComposerClient* client = new SurfaceComposerClient();
```

- SurfaceComposerClient is the interface to SurfaceFlinger.
- It's responsible for asking SurfaceFlinger to create a new surface.

7. SurfaceComposerClient::createSurfaceChecked()

```
cpp
sf->createConnection() → new Client() → new SurfaceControl
```

- Now the **real surface** is created by talking to **SurfaceFlinger**.
- Internally, it uses **BufferQueue + GraphicBuffer + SurfaceControl**.

8. Back in App: ViewRootImpl.performTraversals()

```
java
relayoutResult = relayoutWindow(...);
```

• Triggers a layout/draw cycle and relayouts the window with the real size and Surface.

9. relayoutWindow() \rightarrow copies Surface from SurfaceControl

java

mSurface.copyFrom(mSurfaceControl);

- Finally, the app's Surface is filled using SurfaceControl.
- Now the app can draw to this surface using Canvas or OpenGL.

10. Draw Frame → Render to Screen

```
java
performDraw();
```

- The app renders using Canvas, RenderThread, or OpenGL.
- SurfaceFlinger will later compose this surface to the display.

Key Concepts Explained

SurfaceControl

- Think of it as a **handle to a layer** managed by SurfaceFlinger.
- It doesn't draw directly; it **represents metadata** (size, position, z-order).

Surface

- This is the actual memory buffer (double/triple buffer) you draw to.
- The app writes pixels to it via Canvas, HardwareRenderer, or OpenGL.

SurfaceSession

- A container for managing multiple SurfaceControls.
- Created only once per window/session.

SurfaceFlinger

- Android's system compositor.
- It gets Surfaces from all apps, composes them into one image, and sends it to the display.

Visualization



```
Activity.makeVisible()
WindowManager.addView()
ViewRootImpl created → setView()
mWindowSession.addToDisplay() → Binder IPC
WMS.addWindow() → creates SurfaceSession
SurfaceComposerClient (via JNI)
SurfaceControl created by SurfaceFlinger
Surface.copyFrom(SurfaceControl)
App can now draw → SurfaceFlinger displays it
```

Summary

- Your **Activity window** ends up as a **Surface** inside **SurfaceFlinger**.
- The WMS mediates the window's lifecycle.
- SurfaceControl is a controller.
- **Surface** is what your app draws on.
- **SurfaceFlinger** is what shows it on screen.

Let me know if you want the **drawing** side (Choreographer + RenderThread + VSync), or how SurfaceFlinger composes all layers.



yes

Great! Let's now walk through the **drawing side** of the pipeline in Android — how your app actually renders to the screen **after** the Surface has been created.

Overview: From requestLayout() to Pixels on Screen

After the **Surface** is created, the actual **drawing** goes through:

```
mathematica
App Thread → ViewRootImpl → Choreographer → RenderThread → SurfaceFlinger → Display
```

🧱 Drawing Pipeline: Step-by-Step

1. ViewRootImpl.requestLayout() triggers traversal

```
java
requestLayout(); → scheduleTraversals(); → doTraversal(); → performTraversals();
```

- This is initiated when the UI changes (like TextView.setText()).
- Schedules a frame to render using Choreographer.

2. Choreographer schedules a vsync callback

```
java
mChoreographer.postCallback(CALLBACK_TRAVERSAL, mTraversalRunnable, null);
```

- Choreographer is a **frame scheduler** tied to display vsync (typically 60fps or 120fps).
- It posts TraversalRunnable that ends up calling performTraversals().

3. ViewRootImpl.performTraversals() runs layout, measure, draw

```
java
view.measure(...) → view.layout(...) → view.draw(canvas)
```

- Calls all the View tree methods to measure and layout.
- Then calls draw(Canvas) which draws everything onto a Canvas.

4. How does Canvas draw to the screen?

There are **two paths**:

A. Software rendering:

- Uses **Skia** and renders to a **Bitmap** backed by the **Surface**.
- Used in safe mode or non-hardware accelerated Views.

B. Hardware acceleration (most apps):

Uses RenderThread + OpenGL ES or Vulkan.



• Backed by HardwareRenderer and ThreadedRenderer.

java

mAttachInfo.mThreadedRenderer.draw(...) → RenderThread → EGL → Surface

5. RenderThread prepares the frame (Hardware-accelerated apps)

- RenderThread runs in parallel to the main thread.
- Uses **OpenGL or Vulkan** to render Views into buffers attached to the Surface.
- It calls **EGL draw calls** into a framebuffer (double/triple buffer).

6. Final frame is posted to SurfaceFlinger

- eglSwapBuffers() posts the final frame to Surface.
- Now **SurfaceFlinger** gets the buffer from your Surface.

What SurfaceFlinger Does

SurfaceFlinger collects surfaces from:

- Apps
- System UI (status bar, navigation bar)
- Input layers
- Wallpaper, etc.

Then it:

- 1. Composites all layers
- 2. **Applies transformations** (rotation, scaling, etc.)
- 3. Uses HWC (Hardware Composer HAL) to push to display

Example Timeline (Simplified at 60fps)

| Time (ms) | Component | Description |
|-----------|--------------|--|
| 0 ms | Арр | Calls requestLayout() |
| 1 ms | ViewRootImpl | Schedules frame via Choreographer |



| Time (ms) | Component | Description |
|-----------|----------------|--|
| 16 ms | VSync | Choreographer.doFrame() triggers traversal |
| 17–22 ms | Layout & Draw | Views drawn to Surface (Canvas or GPU) |
| 23 ms | EGL | eglSwapBuffers() posts buffer |
| 24–32 ms | SurfaceFlinger | Composites and shows on screen |

Visualization

```
App code

i requestLayout()

Choreographer schedules VSync

View hierarchy measured + laid out + drawn

HardwareRenderer or Canvas draws to Surface

SurfaceFlinger composites Surfaces

Display shows final frame
```

Key Components Recap

| Component | Role | |
|---|---|--|
| ViewRootImpl | Bridge between View system and Surface | |
| Choreographer | Schedules frame on vsync | |
| RenderThread | Offloads rendering from UI thread (OpenGL/Vulkan) | |
| SurfaceFlinger | Composes final frame from all surfaces | |
| SurfaceControl | Low-level wrapper to manage a layer | |
| Surface | Actual drawing buffer | |
| EGL | Interface to GPU for rendering | |
| HWC Talks to display hardware (via HAL) | | |

+ Want to Go Deeper?

I can explain any of these in detail:

- RenderThread architecture and EGL setup
- SurfaceFlinger composition pipeline



- How VSync and triple buffering reduce jank
- Difference between SurfaceView and TextureView
- GPU memory and buffer management (GraphicBuffer, BufferQueue)
- Hardware Composer HAL (composer@2.x)

Just tell me which one you want to dive into.



Surface display process summary:

After creating PhoneWindow in the App process, ViewRoot is created. The creation of ViewRoot creates a Surface, which is actually empty. When communicating with WindowManagerService, copyFrom() a NativeSurface is used to communicate with SurfaceFlinger.

About Native Window:

Native Window is a bridge between OpenGL and the local window system. The entire GUI system requires at least two types of local windows:

(1) Manager-oriented (SurfaceFlinger)

SurfaceFlinger is the manager of all UI interfaces in the system and needs to directly or indirectly hold a "local window". This local window is FramebufferNativeWindow.

(2) Application-oriented

This type of local window is Surface.

Normally, when generating an APK application according to the SDK wizard, third-party graphics libraries such as Skia are used. However, for application developers who want to use OpenGL ES to complete complex interface rendering, Android also provides encapsulated GLSurfaceView (or other methods) to achieve graphics display.

(1)FramebufferNativeWindow

EGL needs to use a native window to create an environment for OpenGL/OpenGL ES. Since OpenGL/ES supports multiple platforms, compatibility and portability are taken into consideration. The data types of local windows EGLNativeWindowType on different platforms are different.

The data type on the Android platform is ANativeWindow, which is like a "protocol" that specifies the form and function of a local window. ANativeWindow is the parent class of FramebufferNativeWindow.

In Android, due to multi-buffer technology, the buffers managed by



EGLNativeWindowType are at least 2 and at most 3.

FramebufferNativeWindow initialization requires Golloc support, and the steps are as follows:

Load the GRALLOC_HARDWARE_MODULE_ID module, see the previous section.

Open the fb and gralloc devices respectively. The opened devices are managed by the global variables fbDev and grDev.

Assign initial values to FramebufferNativeWindow according to the properties of the device. Fill in the "protocol" in ANativeWindow according to the implementation of FramebufferNativeWindow

Some other necessary initialization

②The local window of the application - Surface Surface also inherits ANativeWindow

class Surface: public ANativeObjectBase<ANativeWindow, Surface, RefBase>

Surface is the interface for all UI applications in the Android system, fulfilling the UI display requirements of the application process.

It provides a drawing board for upper-layer implementations (primarily the Java layer) to draw images. SurfaceFlinger collects image data drawn by all applications in the system and displays it centrally on the physical screen. The Surface, while fulfilling this role, is essentially managed by the SurfaceFlinger service, which involves many cross-process communication details.

③Surface Creation:

The Surface obtains buffers through mGraphicBufferProducer, which are recorded in the data in mSlots. The initialization process for this core component, mGraphicBufferProducer, is as follows:

ViewRootImpl holds a Java-level Surface object (mSurface).

ViewRootImpl initiates a relayout request to WindowManagerService, at which point mSurface is assigned a valid value. The resulting SurfaceControl is then copied to mSurface via the Surface.copyFrom() function.

Thus, the Surface is managed by the SurfaceControl, which is created by the SurfaceComposerClient. The anonymous Binder obtained by the SurfaceComposerClient is ISurfaceComposer, whose server-side implementation is SurfaceFlinger . The IGraphicBufferProducer object that the Surface relies on is implemented as a BufferQueue on the server side.

class SurfaceFlinger:

public BinderService<SurfaceFlinger>, // Register as SurfaceFlinger in ServiceManager public BnSurfaceComposer, // The implemented interface is called ISurfaceComposer

4 SurfaceFlinger Service Framework:

Buffer, Consumer, and Producer are the three participants in the "Producer-Consumer" model. How to coordinate their work is key to whether the application can display the UI properly.



The buffer is BufferQueue, the Producer is the application, and the Consumer is SurfaceFlinger.

Great summary! Here's a **concise and clearer breakdown** of the entire **Surface display process** and **native window usage in Android**, tying everything together in a digestible way:



1. Surface Creation and Display Flow

High-level steps:

- 1. **App starts Activity** → Activity.makeVisible()
- 2. PhoneWindow is created.
- 3. ViewRootImpl is created \rightarrow owns a Java-level Surface.
- 4. ViewRootImpl.relayoutWindow() calls WindowManagerService (WMS).
- 5. WMS requests a SurfaceControl from SurfaceFlinger.
- 6. SurfaceControl is wrapped into a **native Surface** via **Surface.copyFrom()**.
- 7. App renders into the **Surface** via Skia, Canvas, or OpenGL (via EGL).
- 8. **SurfaceFlinger** composites all surfaces and sends final frame to display.

2. What is a Layer?

- A Layer = Surface + SurfaceControl
- It represents a unit of composition that SurfaceFlinger can manage.
- Think of layers like transparent sheets stacked on top of each other each app's surface is one layer.

3. Understanding Native Window Types

ANativeWindow

- A **platform abstraction** used by EGL/OpenGL ES for rendering targets.
- Implemented as either Surface or FramebufferNativeWindow.

✓ Application-side: Surface

- Inherits from ANativeWindow
- Used by:



- Skia (Java Canvas)
- EGL/OpenGL via EGLWindowSurface
- MediaCodec output
- Behind it is a BufferQueue:
 - **Producer:** App (draws to buffer)
 - **Consumer:** SurfaceFlinger (composites it)

System-side: FramebufferNativeWindow

- Also inherits ANativeWindow
- Used by SurfaceFlinger to directly render to the physical framebuffer (when HWComposer is not available).
- Usually involved in boot logo / recovery modes.

4. SurfaceFlinger Role

- SurfaceFlinger is the **central compositor** in Android.
- It:
- Takes all app surfaces (each with a BufferQueue).
- Reads their contents (consumer side).
- Composites them using hardware composer (HWC) or GPU.
- Displays final output on screen.

5. SurfaceControl & SurfaceComposerClient

- SurfaceComposerClient (in app or WMS) requests SurfaceControl from SurfaceFlinger.
- SurfaceControl creates/manages a layer.
- That control creates a **BufferQueue** and gives the producer end to the app (**Surface**).
- The consumer end is managed by SurfaceFlinger.

6. BufferQueue in Detail

• A BufferQueue is the heart of the producer-consumer mechanism.

| Component | Role | Notes |
|---------------|----------|--------------|
| App / Surface | Producer | Draws frames |

| Component | Role | Notes |
|----------------|----------|----------------------------|
| SurfaceFlinger | Consumer | Reads frames for display |
| BufferQueue | Shared | Coordinates buffer handoff |

- It maintains at least 2 or 3 buffers (triple buffering).
- Uses gralloc to allocate buffers in shared memory (ashmem or dmabuf-backed).

7. EGL Integration

• EGL needs a native_window (i.e., ANativeWindow) to create a rendering surface:

```
c
EGLSurface eglSurface = eglCreateWindowSurface(eglDisplay, config, native_window, NULL);
```

• If you're using OpenGL ES in a game or GL-based app, you're rendering into a **Surface** (ANativeWindow) backed by a **BufferQueue**.

> 8. Simple Real-World Example

Let's say you open YouTube:

- 1. The app creates a Window and Surface.
- 2. Video rendering uses **MediaCodec** → **Surface**.
- 3. UI (play button, seek bar) is drawn on the same or another surface.
- 4. SurfaceFlinger:
 - Gets these surfaces.
 - Composites the video + UI layers.
 - Shows the result on screen.

🔚 Summary Diagram (simplified):

```
App: ViewRootImpl

Surface (ANativeWindow) --> Producer ---

SurfaceFlinger

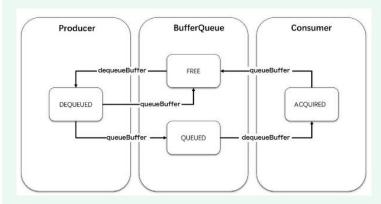
BufferQueue

Framebuffer / DisplayController <-- Consumer
```



Let me know if you want this broken down visually with a diagram or want to dive into a specific part (e.g., SurfaceControl, gralloc, BufferQueue, etc.)





Surface provides a BufferQueue internally, which forms a producer-consumer model with the upper layer and SurfaceFlinger. The upper layer corresponds to the Producer, and SurfaceFlinger corresponds to the Consumer. The three are connected through the Buffer, and each Buffer has four states:

Free: can be used by upper layers;

Dequeued: out of the gueue and being used by the upper layer;

Queued: Enqueued, upper layer drawing has been completed, waiting for SurfaceFlinger synthesis;

Acquired: Acquired, SurfaceFlinger is holding the Buffer for synthesis;

This cycle forms a process in which the Buffer is recycled (FREE->DEQUEUED->QUEUED->ACQUIRED->FREE).

The mSlots array in BufferQueue is used to manage the buffer in the period, and the maximum container is 32. The space of the data buffer is dynamically allocated. Both the application and SurfaceFlinger use OpenGL ES to complete the UI display. The Layer class represents a "layer" in SurfaceFlinger, which generally represents a "screen". The final display result on the physical screen is formed by processing and superimposing all the "screens" that exist in the system at the same time.

Excellent explanation — and the **diagram** you've shared clearly visualizes the **BufferQueue lifecycle** and its **producer-consumer** mechanics.

Let's walk through the **full cycle of buffer states** and how each component interacts:

Buffer Lifecycle in BufferQueue (via Surface)



1. FREE

- The buffer is available for the producer (usually the app) to use.
- Stored in a pool inside BufferQueue.

2. **DEQUEUED**

- The app (producer) calls dequeueBuffer() to obtain a buffer.
- The buffer is now *locked for writing*.
- App performs rendering (Canvas, OpenGL, MediaCodec, etc.).

3. QUEUED

- After drawing, the app calls queueBuffer() to submit it.
- The buffer enters the *queue*, waiting to be processed by the consumer.

4. ACQUIRED

- SurfaceFlinger (the consumer) calls dequeueBuffer() to obtain the ready buffer.
- It performs compositing blending it with other layers (via HWC/GPU).

5. FREE (again)

• Once SurfaceFlinger is done rendering the frame, it calls queueBuffer() to return the buffer to the free pool.

ℰ Component Roles

| Component | Role | Key Method Calls |
|---------------------------|-------------------|---|
| Producer (App) | Renders UI | $\textbf{dequeueBuffer()} \rightarrow \textbf{draw} \rightarrow \textbf{queueBuffer()}$ |
| BufferQueue | Mediator | Manages slot state transitions |
| Consumer (SurfaceFlinger) | Composites frames | $\textbf{dequeueBuffer()} \rightarrow \texttt{render} \rightarrow \textbf{queueBuffer()}$ |

🔢 Buffer Management: mSlots Array

- BufferQueue maintains up to 32 slots via mSlots[].
- Each slot holds a reference to a buffer and its current **state**.
- Slots are reused to reduce buffer allocation overhead.
- Slots move through state transitions:

FREE → DEQUEUED → QUEUED → ACQUIRED → FREE

Advanced Notes



- Triple Buffering is commonly used:
 - 1 buffer being displayed (ACQUIRED)
 - 1 buffer being drawn into (DEQUEUED)
 - 1 buffer waiting to be displayed (QUEUED)

• Backpressure:

- If the producer is too fast, it may be blocked on dequeueBuffer() until a buffer becomes FREE again.
- Ensures synchronization with display refresh rate (e.g. 60Hz).

• SurfaceFlinger's Layer:

- Each surface/layer in the system is represented by a Layer class in SurfaceFlinger.
- Composited layers form the final screen output.

Simplified Visual Summary

Let me know if you'd like an annotated version of the diagram you uploaded with these state transitions labeled — or a more detailed visual of mSlots[] and buffer slot management.